



Generative Adversarial Network

Rookie Training Camp

Advisor: 李家岩 博士

Presenter: 歐子毓

Date: 2020.07.15

GAN as structured learning

01

Basic idea of GAN

02

Can discriminator generate

03

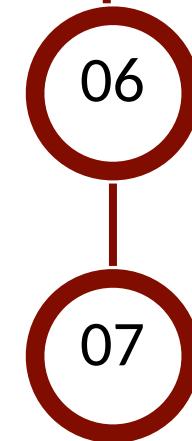
Can generator learn by itself

04

Conditional GAN

05

Conquer some issues using GANS



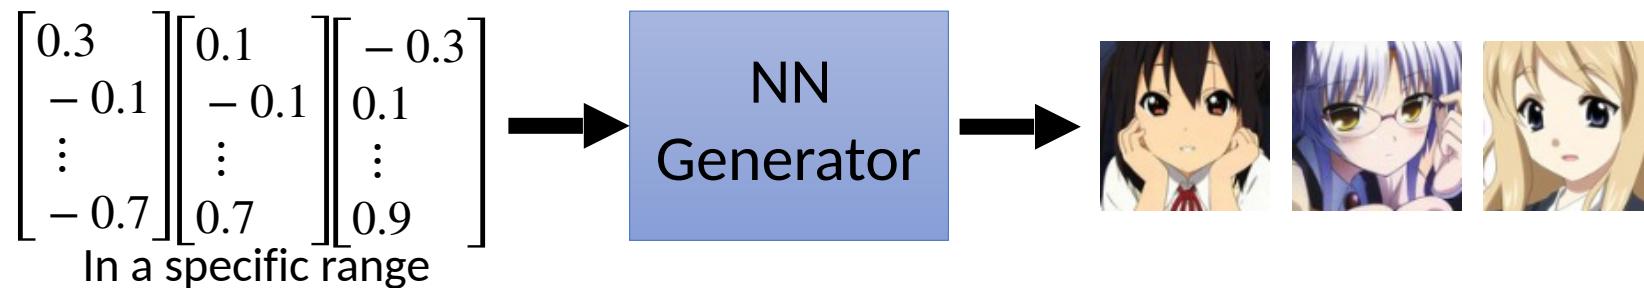
Implementation with PyTorch

Basic idea of GAN

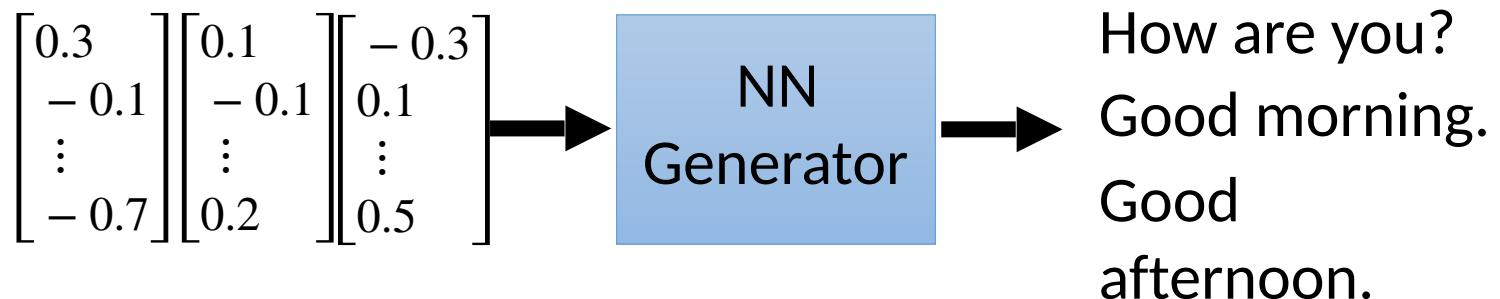
Generation

We will control what to generate latter. → Conditional Generation

Image Generation

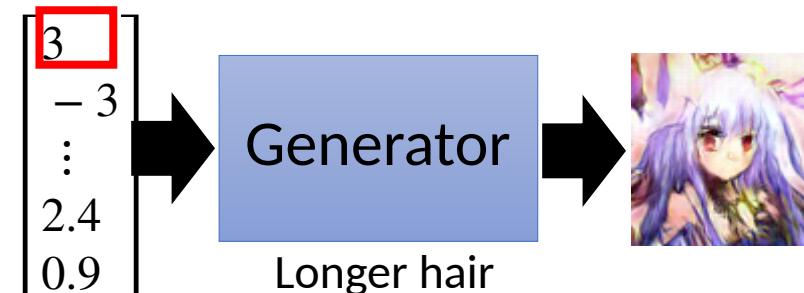
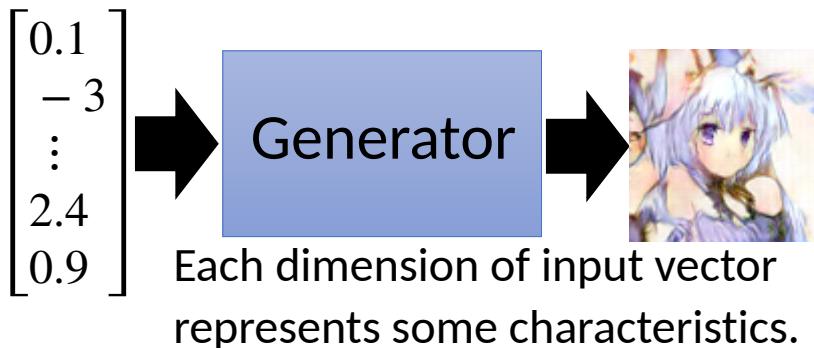
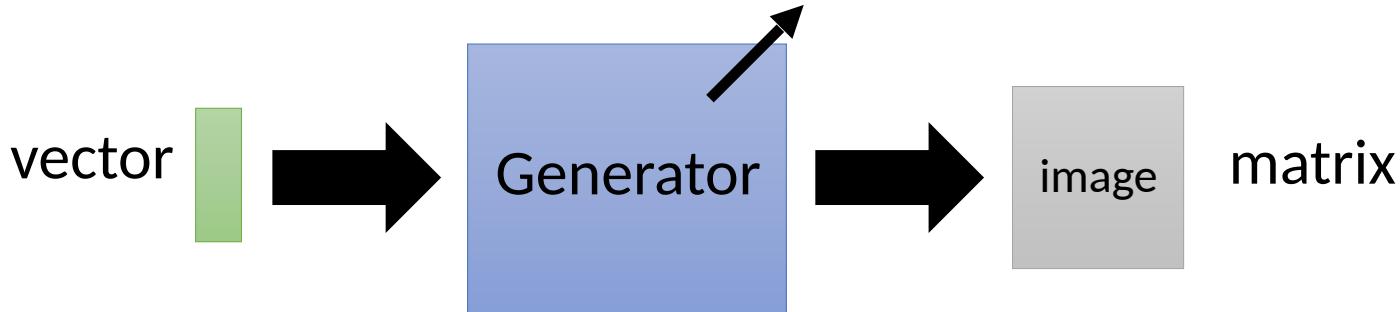


Sentence Generation



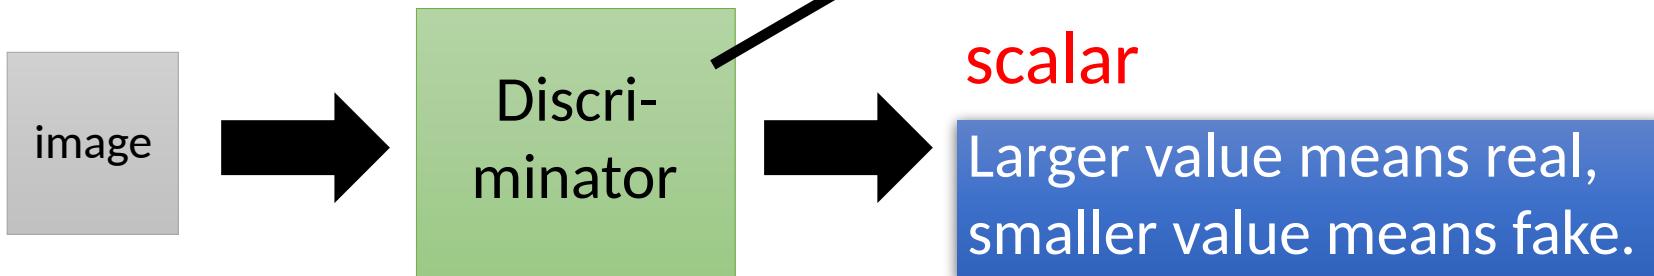
Basic Idea of GAN

It is a neural network (NN), or a function.

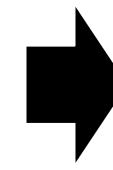
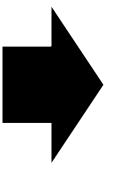


Basic Idea of GAN

It is a neural network (NN), or a function.



Basic Idea of GAN



Generator

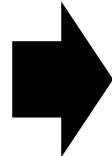
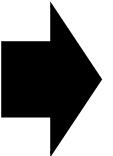
Brown

veins

Butterflies are
not brown

Butterflies do
not have veins

.....

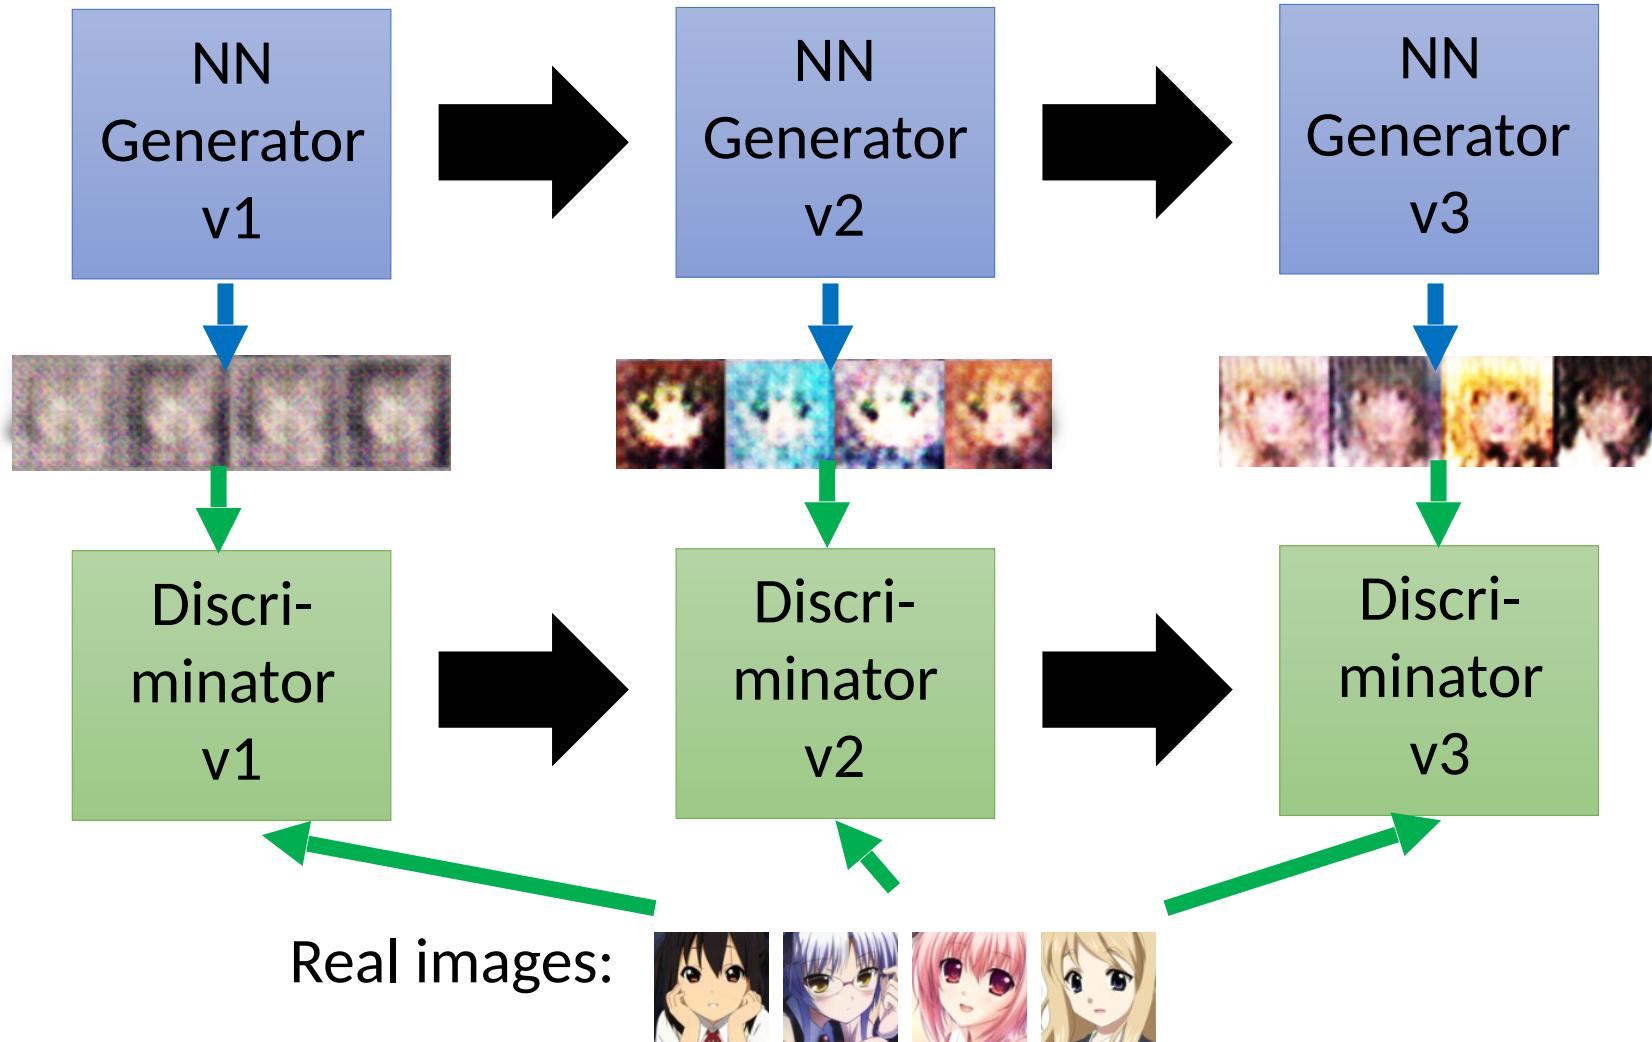


Discriminator

Basic Idea of GAN

This is where the term "*adversarial*" comes from.

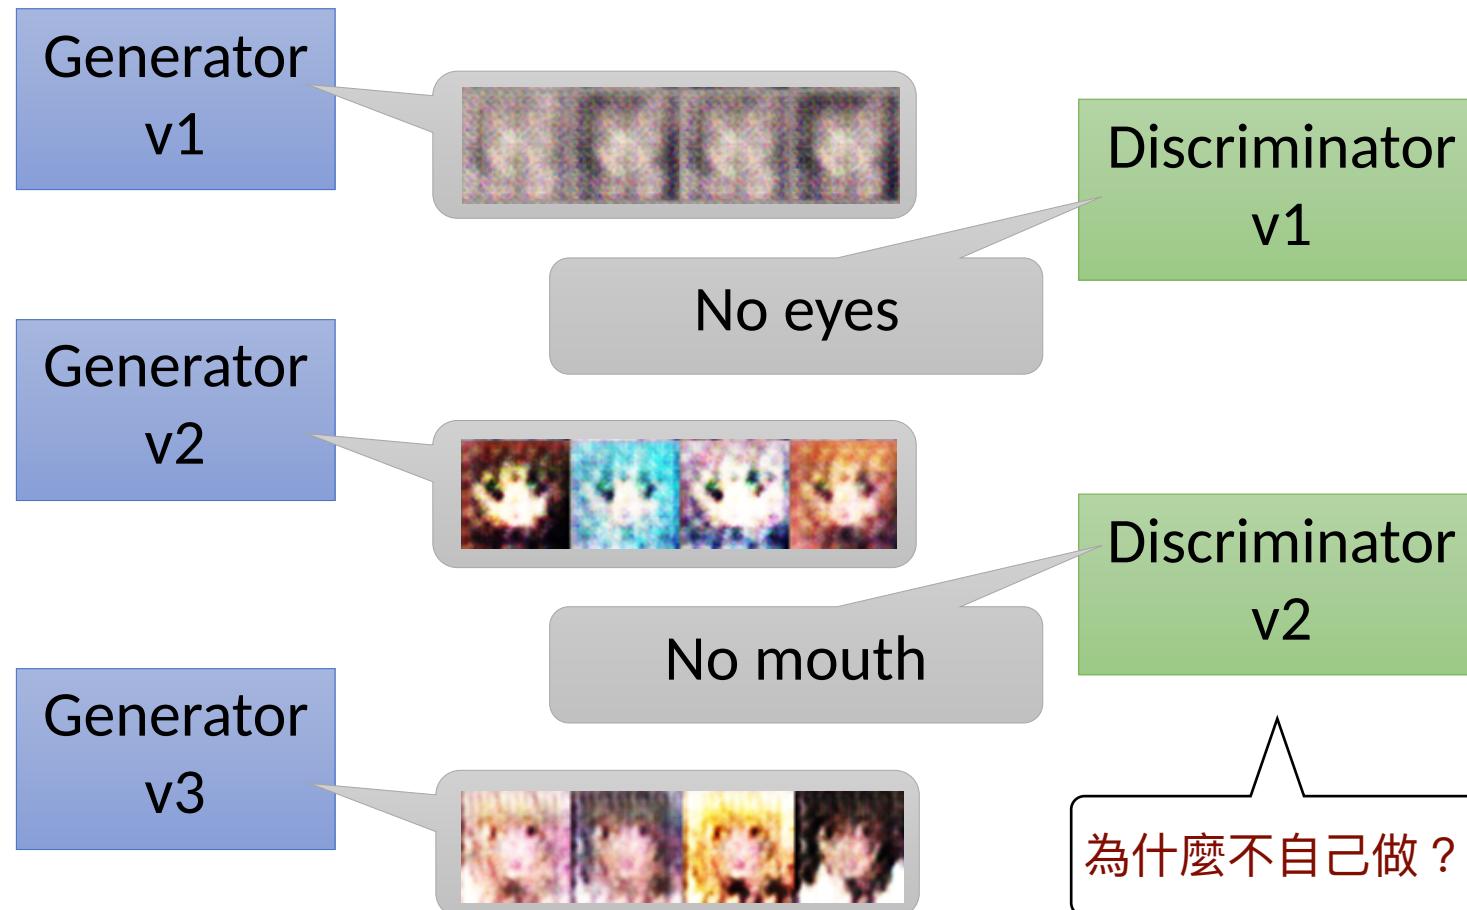
You can explain the process in different ways.....



Basic Idea of GAN

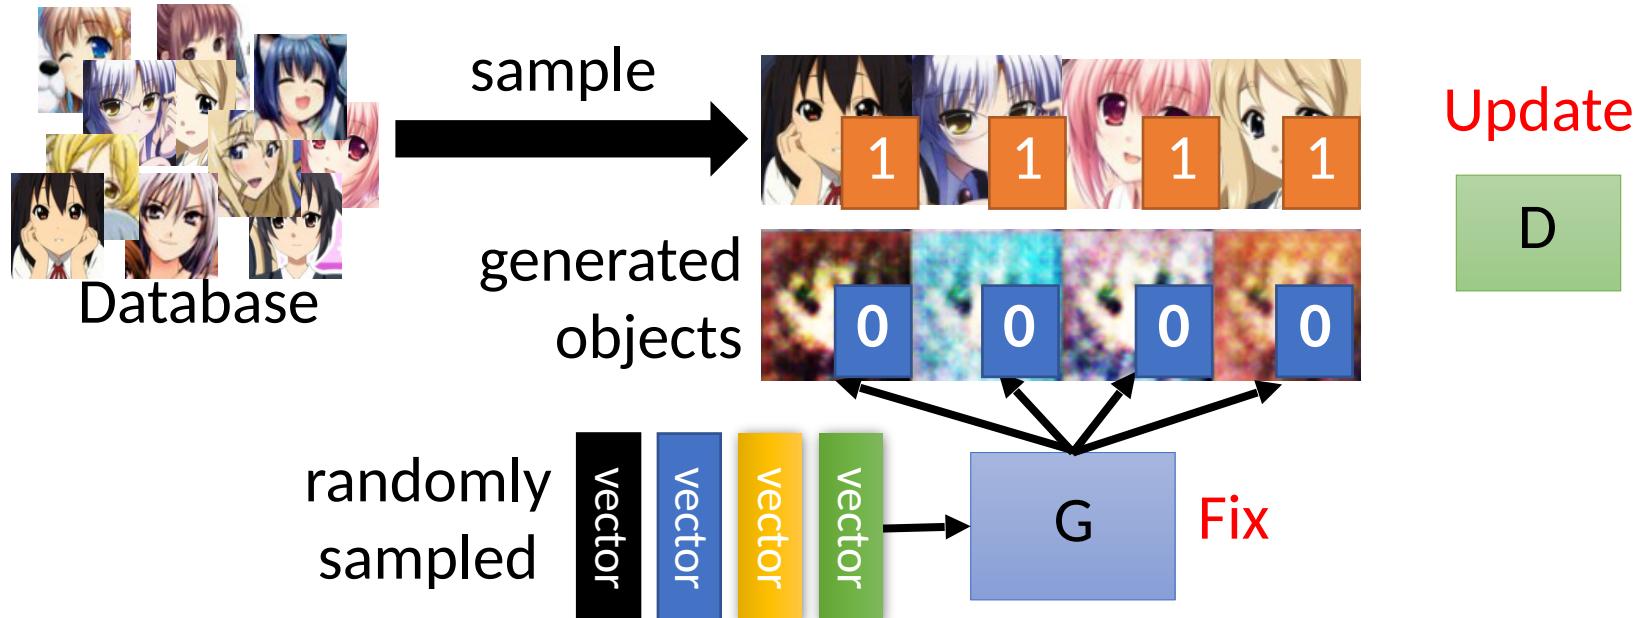
Generator
(student)

Discriminator
(teacher)



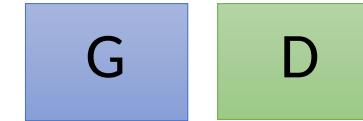
- Initialize generator and discriminator
- In each training iteration:

Step 1: Fix generator G, and update discriminator D



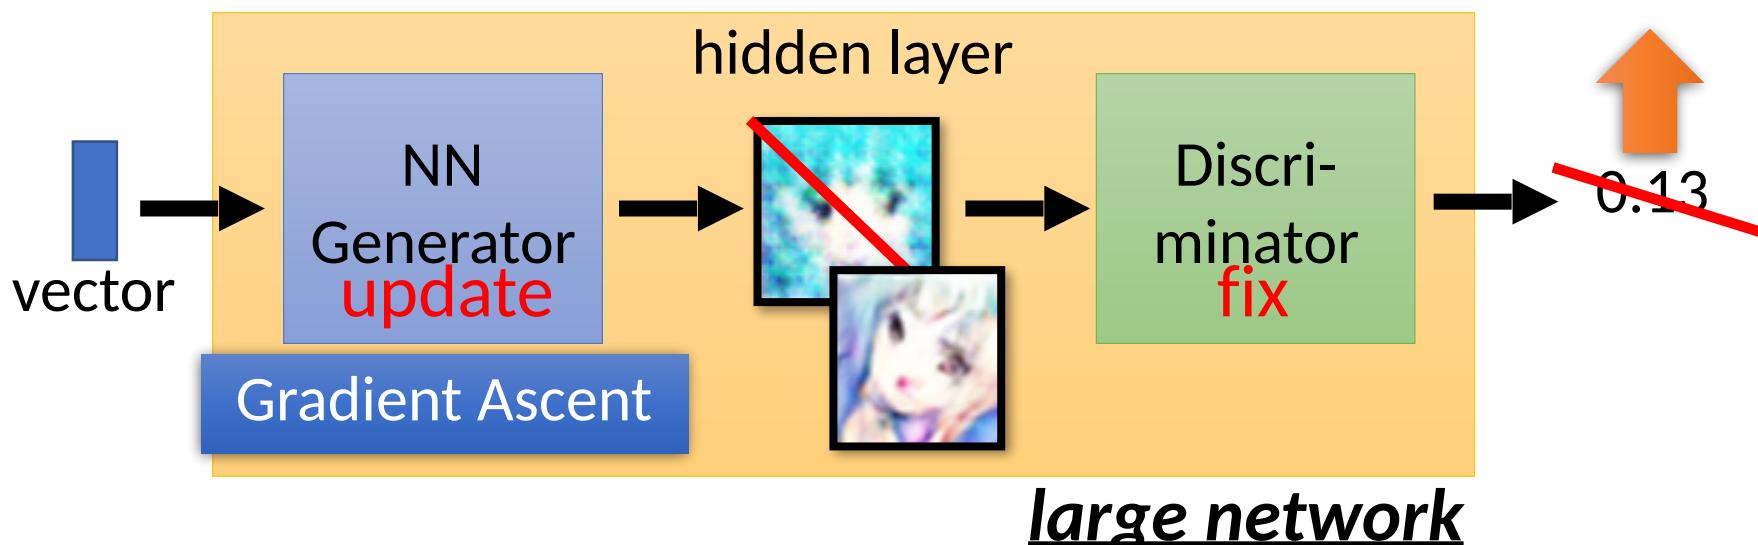
Discriminator learns to assign high scores to real objects and low scores to generated objects.

- Initialize generator and discriminator
- In each training iteration:



Step 2: Fix discriminator D, and update generator G

Generator learns to “fool” the discriminator



Initialize θ_d for D
and θ_g for G

Learning
D

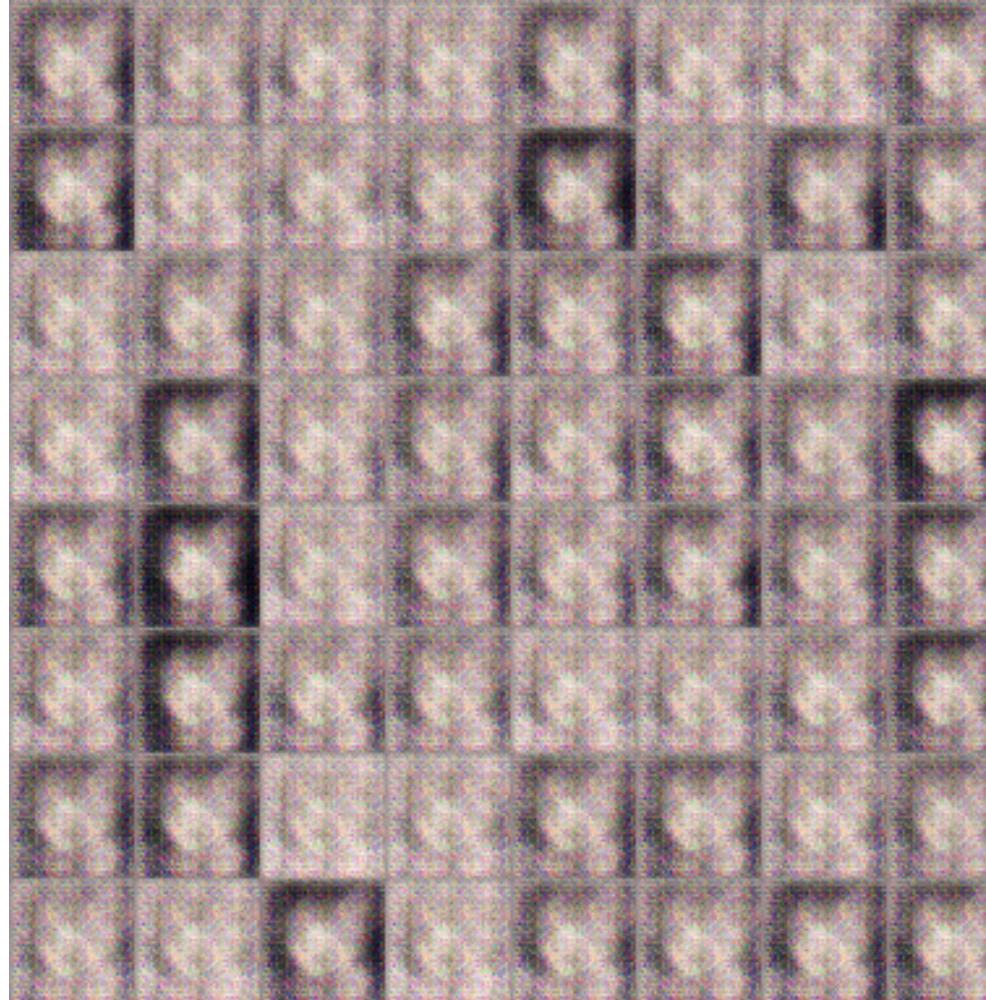
Learning
G

- In each training iteration:

- Sample m examples $\{x^1, x^2, \dots, x^m\}$ from database
- Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from a distribution
- Obtaining generated data $\{\tilde{x}^1, \tilde{x}^2, \dots, \tilde{x}^m\}$, $\tilde{x}^i = G(z^i)$
- Update discriminator parameters θ_d to maximize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log D(x^i) + \frac{1}{m} \sum_{i=1}^m \log(1 - D(\tilde{x}^i))$
 - $\theta_d \leftarrow \theta_d + \eta \nabla \tilde{V}(\theta_d)$

- Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from a distribution
- Update generator parameters θ_g to maximize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log(D(G(z^i)))$
 - $\theta_g \leftarrow \theta_g - \eta \nabla \tilde{V}(\theta_g)$

Anime Face Generation



Source of training data: <https://zhuanlan.zhihu.com/p/24767059>

Anime Face Generation



1000 updates

Anime Face Generation

5000 updates



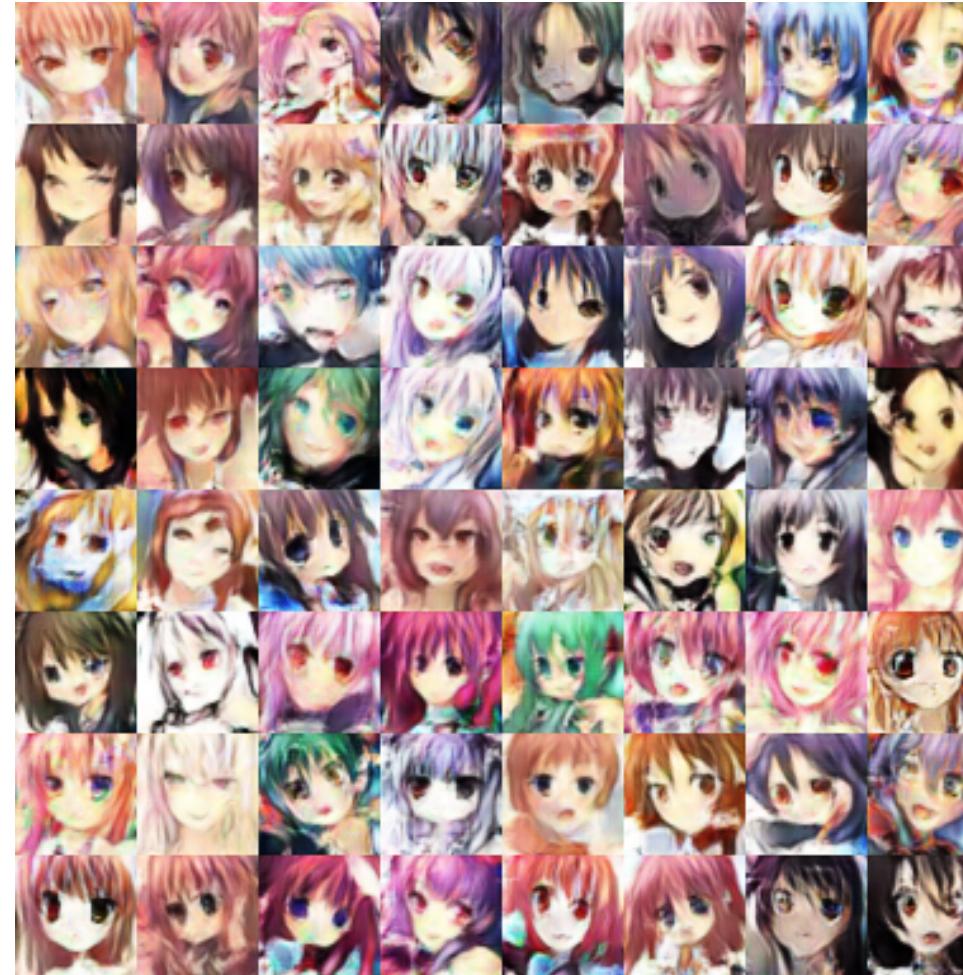
Anime Face Generation

10,000 updates

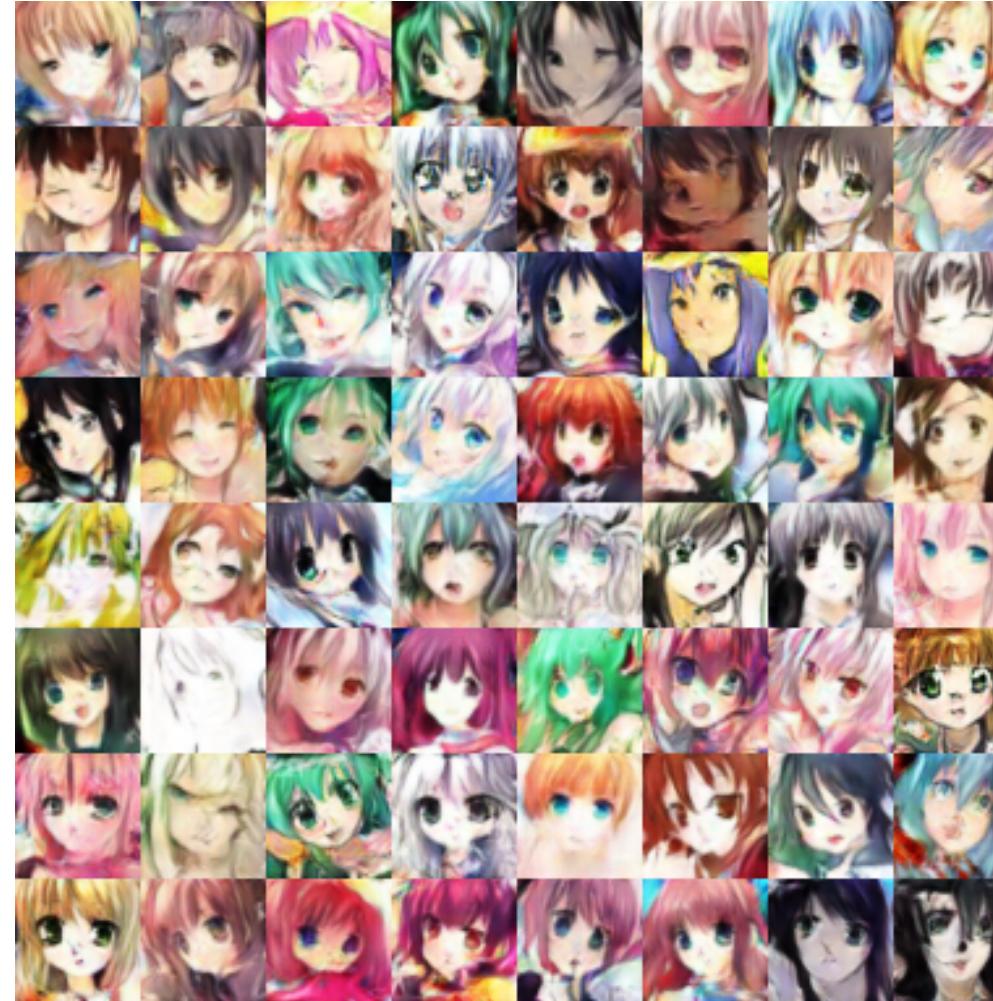


Anime Face Generation

20,000 updates



Anime Face Generation



50,000 updates

GAN as structured learning

Structured Learning

Machine learning is to find a function f

$$f : X \rightarrow Y$$

Regression: output a scalar

Classification: output a “class” (one-hot vector)

		
Class 1	Class 2	Class 3

Structured Learning/Prediction: output a sequence, a matrix, a graph, a tree

Output is composed of components with dependency

Output Sequence

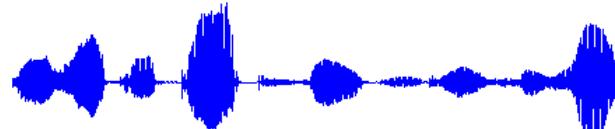
$$f : X \rightarrow Y$$

Machine Translation

X ：“機器學習及其深層與
結構化”
(sentence of language 1)

Y : “Machine learning and
having it deep and
structured”
(sentence of language 2)

Speech Recognition

X : 
(speech)

Y ：感謝大家來上課”
(transcription)

Chat-bot

X ：“How are you?”
(what a user says)

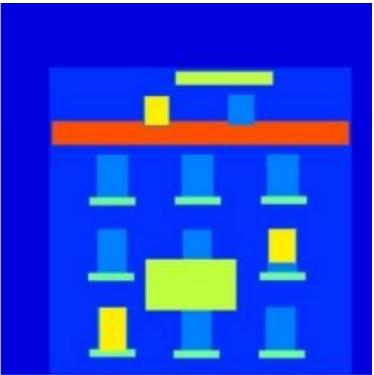
Y ：“I'm fine.”
(response of machine)

Output Matrix

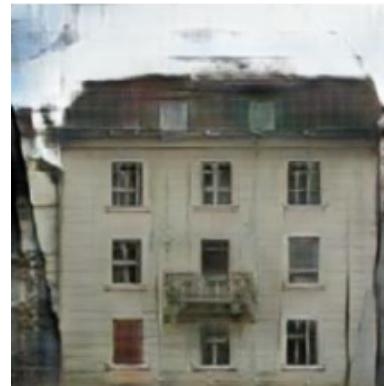
$$f : X \rightarrow Y$$

Image to Image

$X :$



$Y :$



Colorization:



Ref: <https://arxiv.org/pdf/1611.07004v1.pdf>

Text to Image

$X :$ “this white and yellow flower
have thin white petals and a
round yellow stamen”

ref: <https://arxiv.org/pdf/1605.05396.pdf>

$Y :$

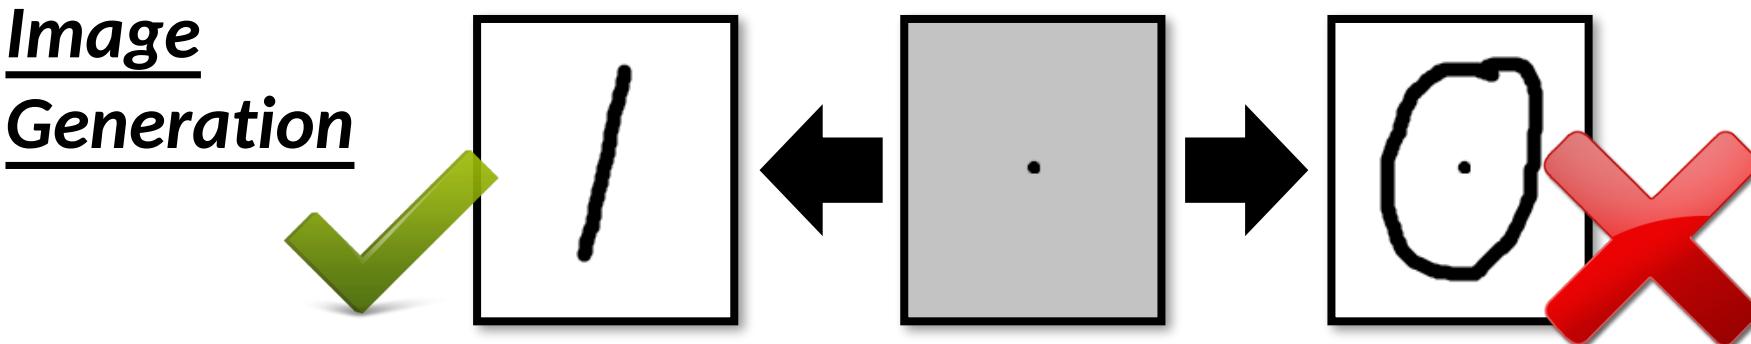


Why Structured Learning Challenging?

- **One-shot/Zero-shot Learning:**
 - In classification, each class has some examples.
 - In structured learning,
 - If you consider each possible output as a “class”
 - Since the output space is huge, most “classes” do not have any training data.
 - Machine has to create new stuff during testing.
 - Need more intelligence

Why Structured Learning Challenging?

- Machine has to learn to do ***planning***
 - Machine generates objects component-by-component, but it should have a big picture in its mind.
 - Because the output components have dependency, they should be considered globally.



Structured Learning Approach

Generator

Learn to generate
the object at the
component level



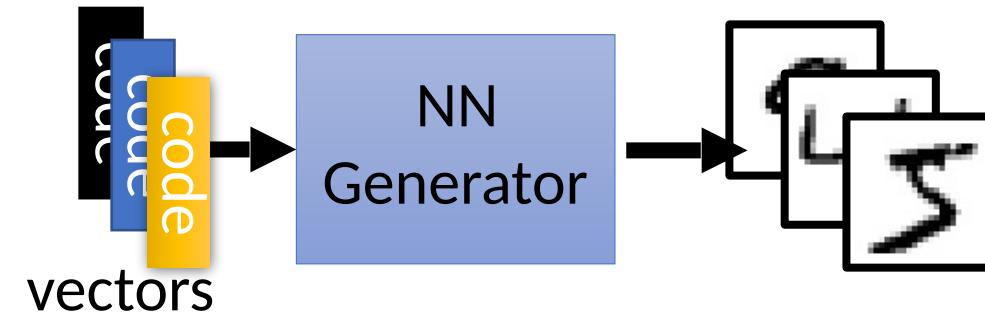
Discriminator

Evaluating the
whole object, and
find the best one



Can generator learn by itself?

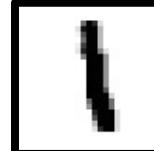
Generator



code:
 (where does they
 come from?)

$\begin{bmatrix} 0.1 \\ -0.5 \end{bmatrix}$

Image:



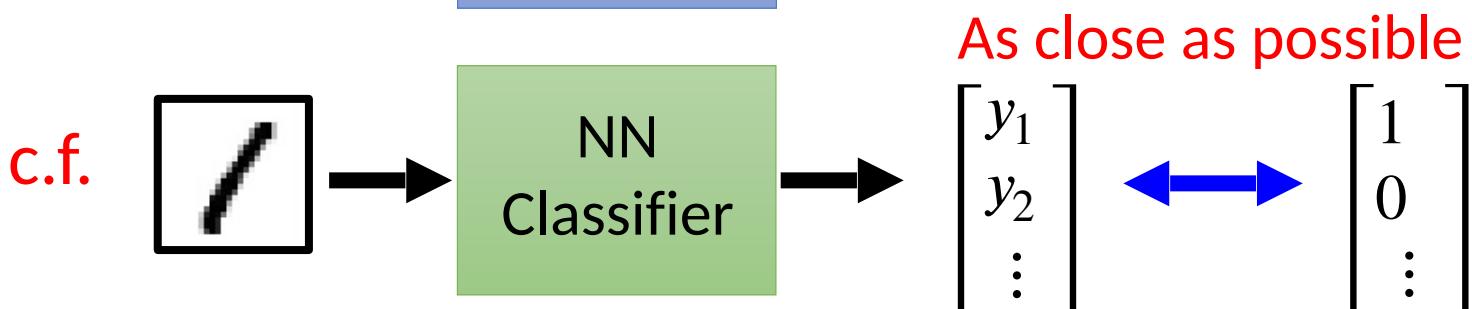
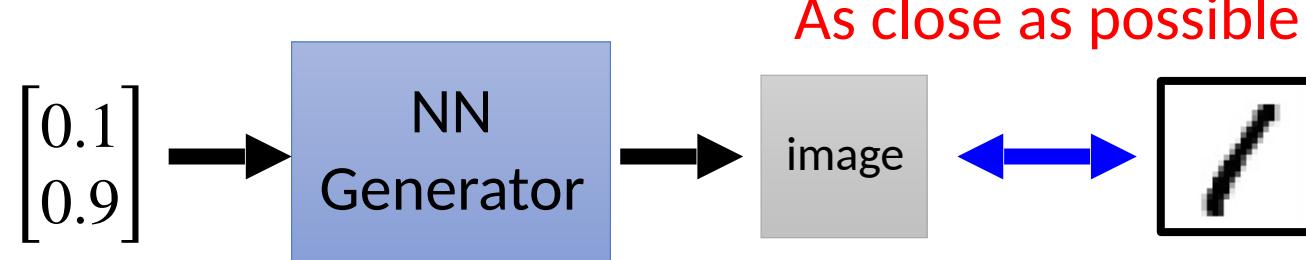
$\begin{bmatrix} 0.1 \\ 0.9 \end{bmatrix}$

$\begin{bmatrix} 0.2 \\ -0.1 \end{bmatrix}$

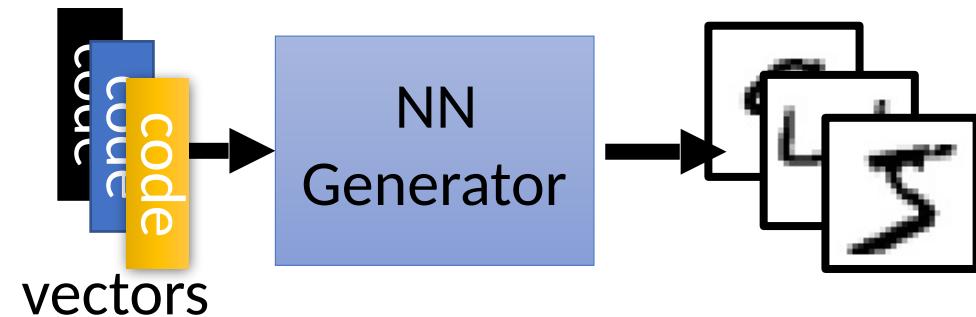
$\begin{bmatrix} 0.3 \\ 0.2 \end{bmatrix}$







Generator

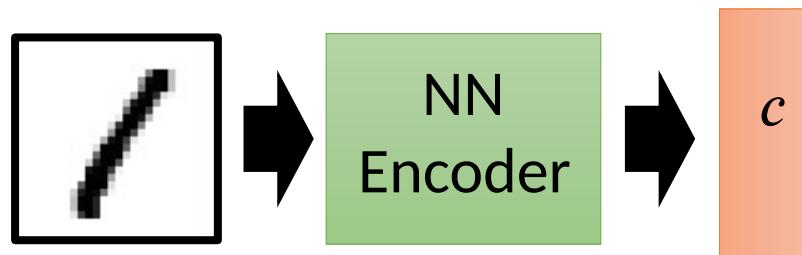


code:
 (where does they
 come from?)

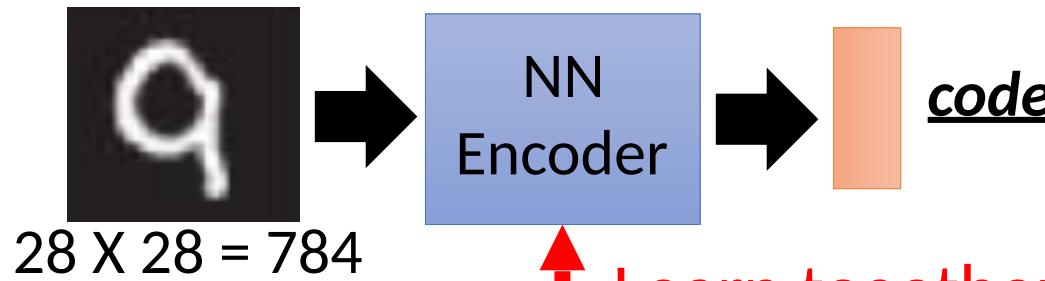
Image:

$\begin{bmatrix} 0.1 \\ -0.5 \end{bmatrix}$	$\begin{bmatrix} 0.1 \\ 0.9 \end{bmatrix}$	$\begin{bmatrix} 0.2 \\ -0.1 \end{bmatrix}$	$\begin{bmatrix} 0.3 \\ 0.2 \end{bmatrix}$

Encoder in auto-encoder
 provides the code ☺



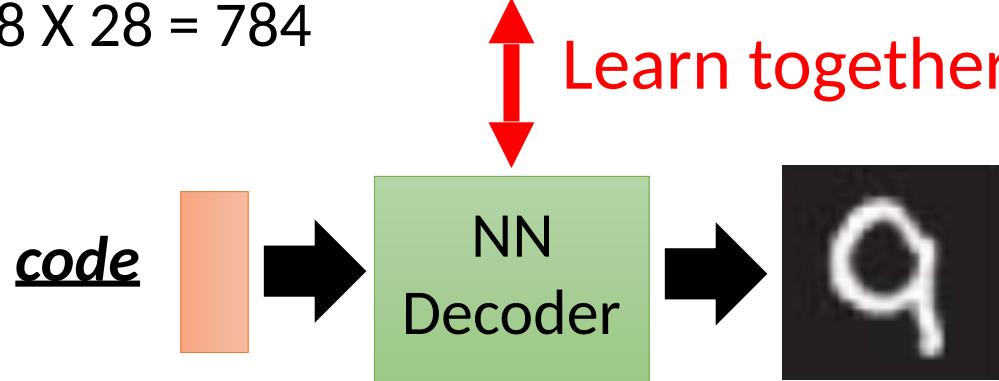
Auto-encoder



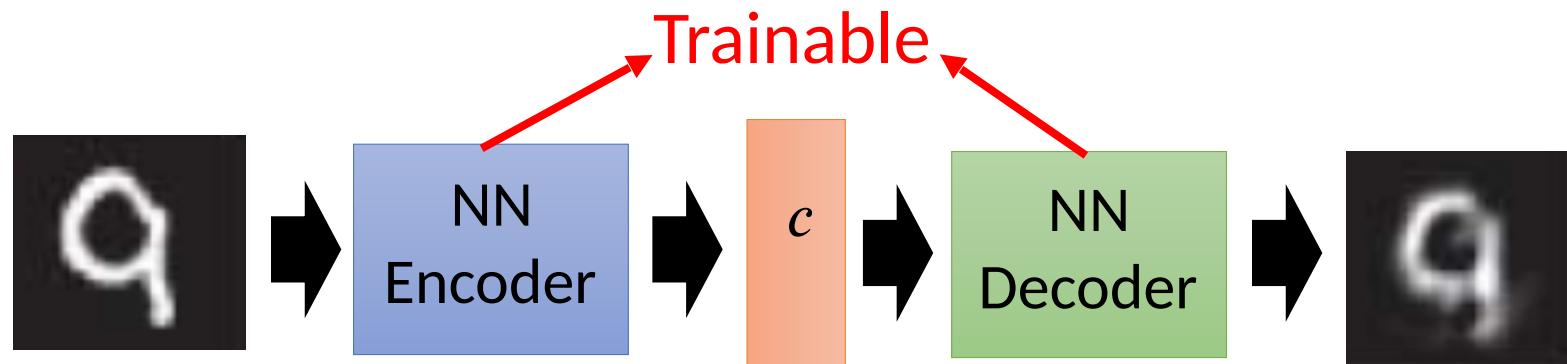
Low dimension

code

Compact representation of the input object



Can reconstruct the original object

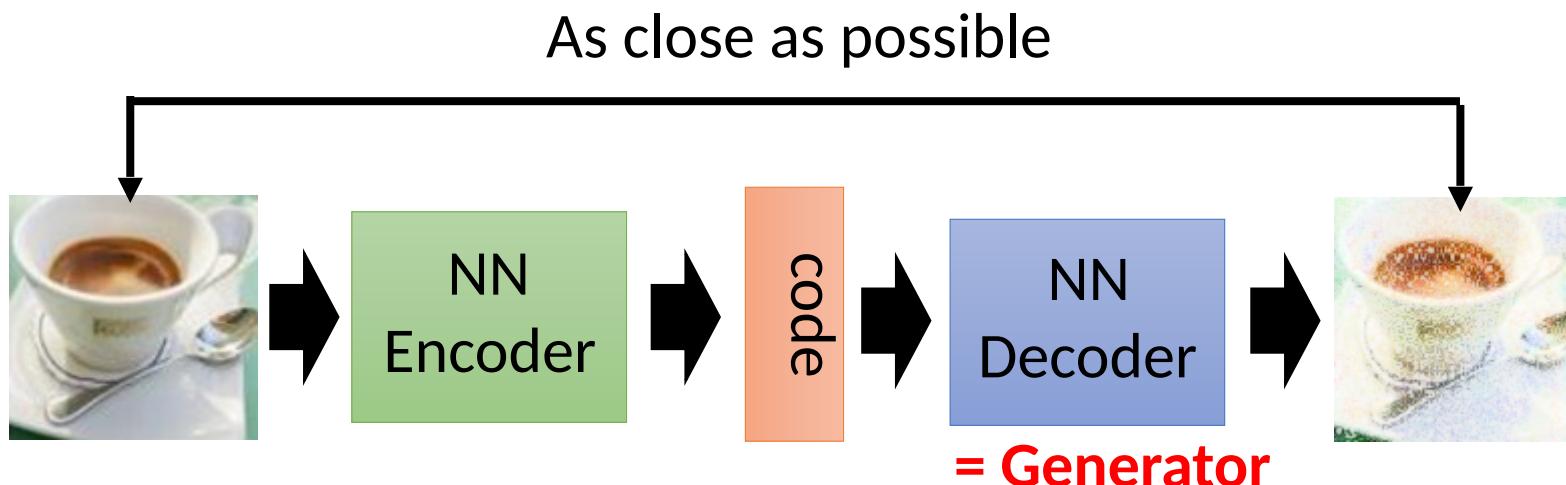


Trainable

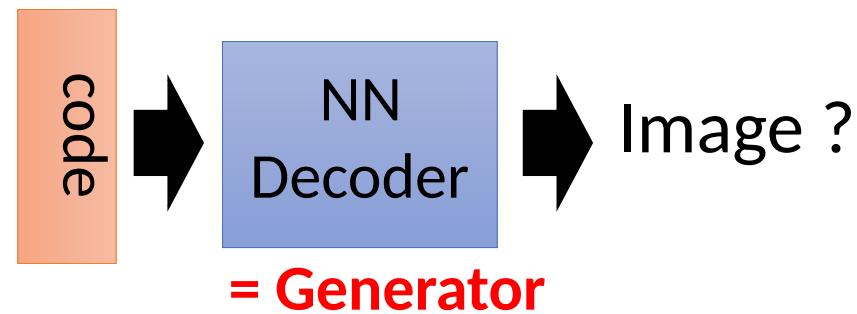
NN
Decoder



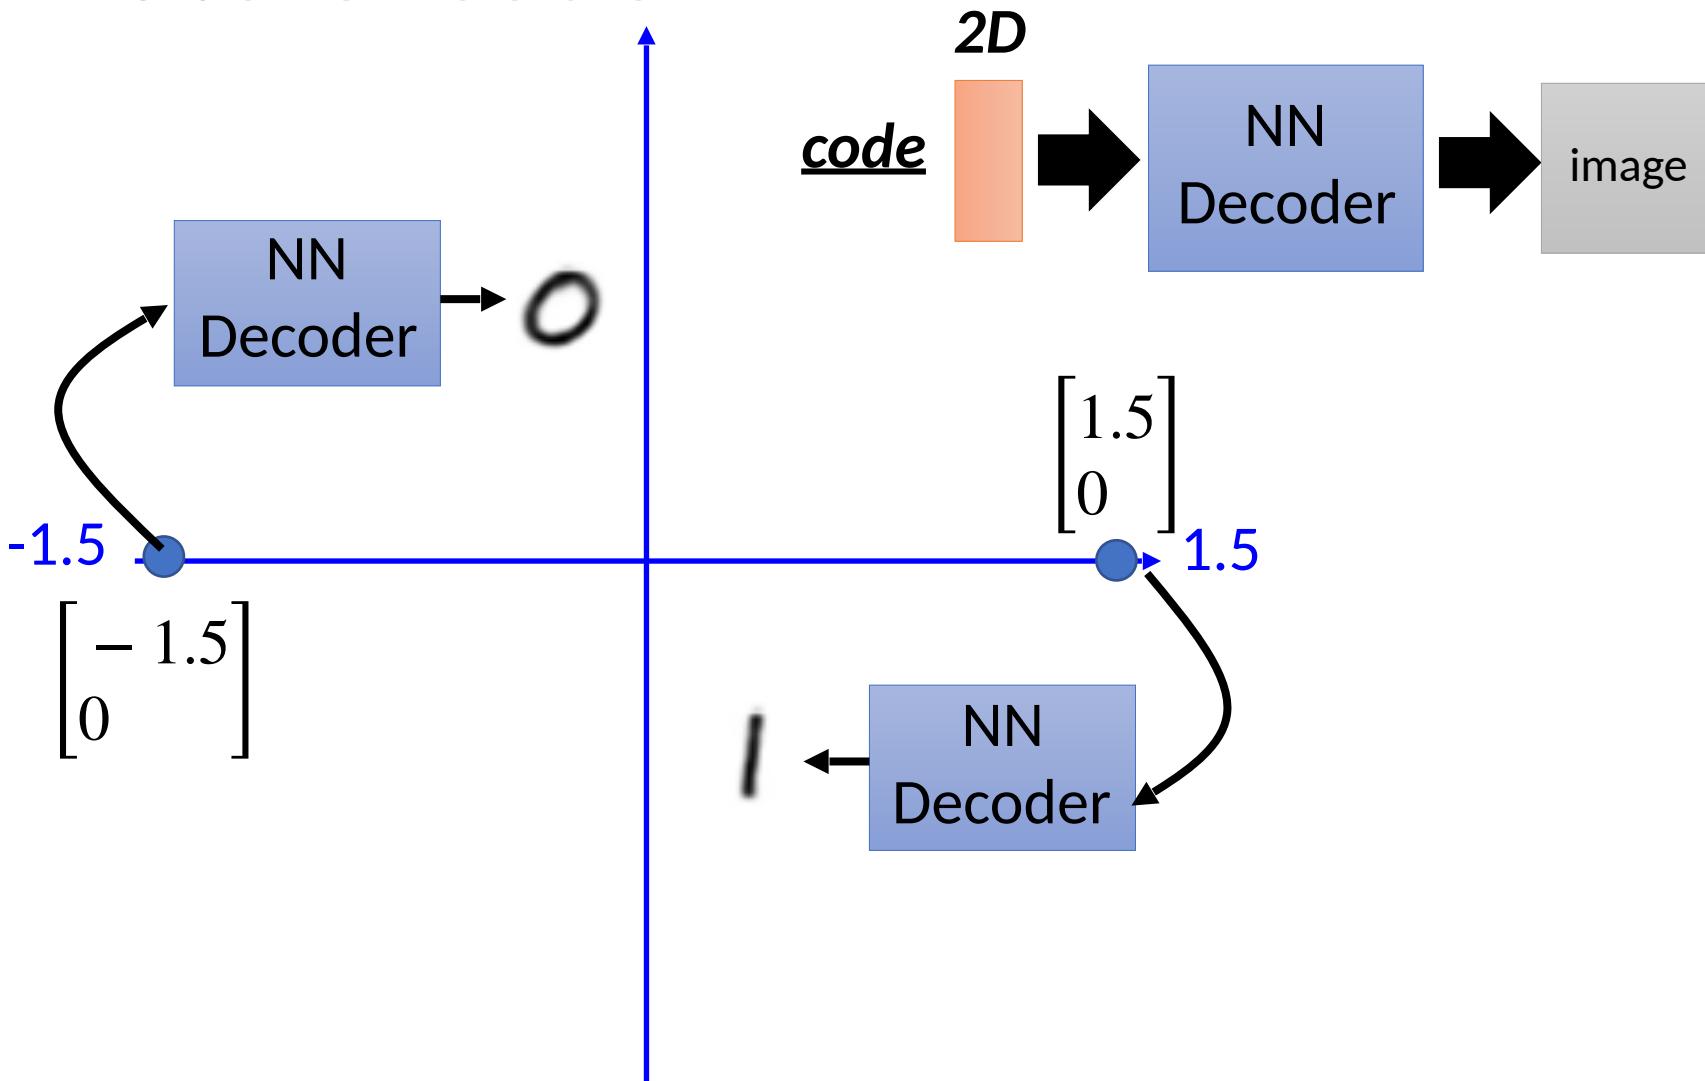
Auto-encoder



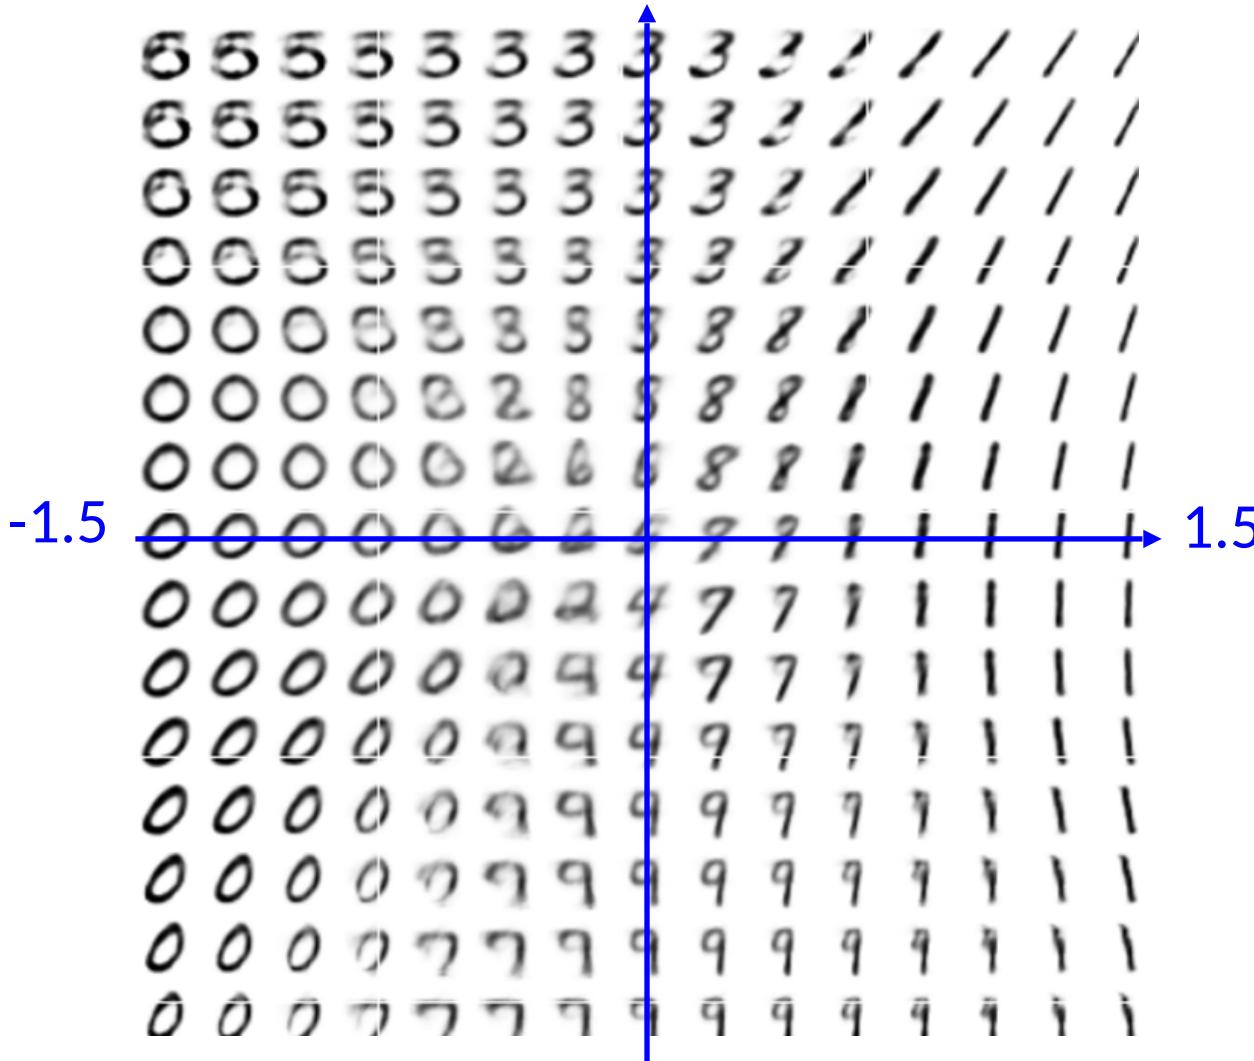
Randomly generate
a vector as code



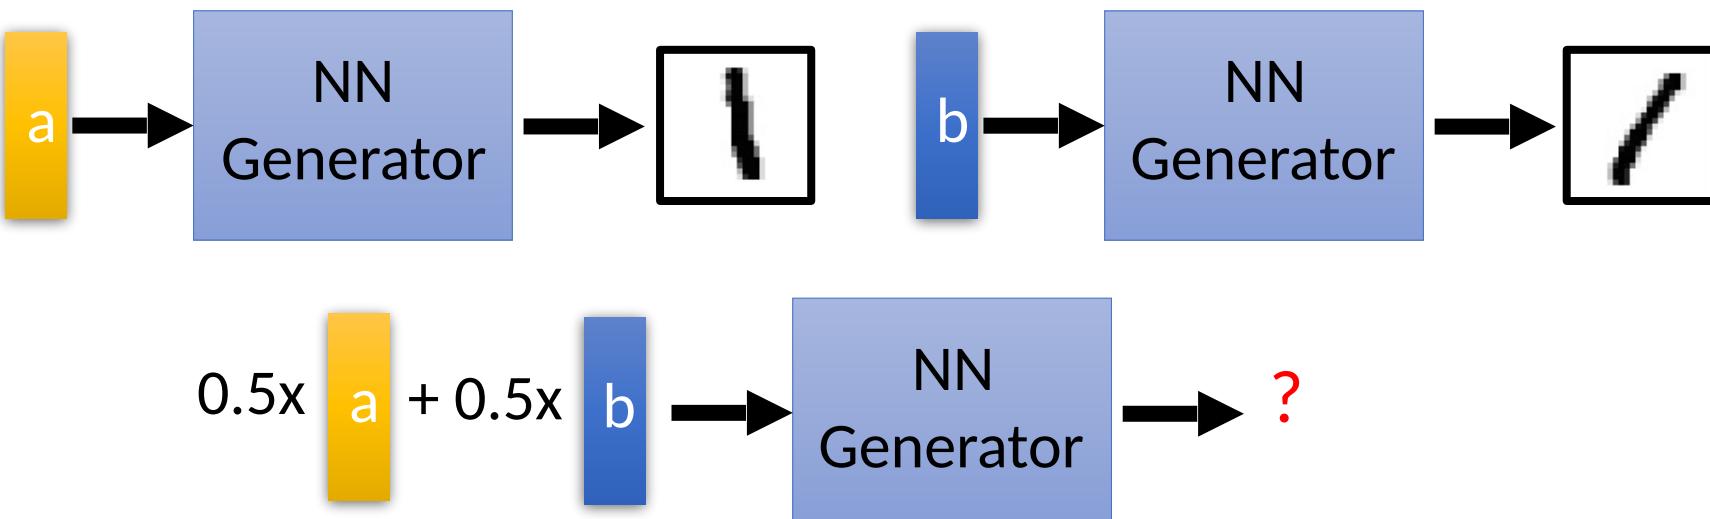
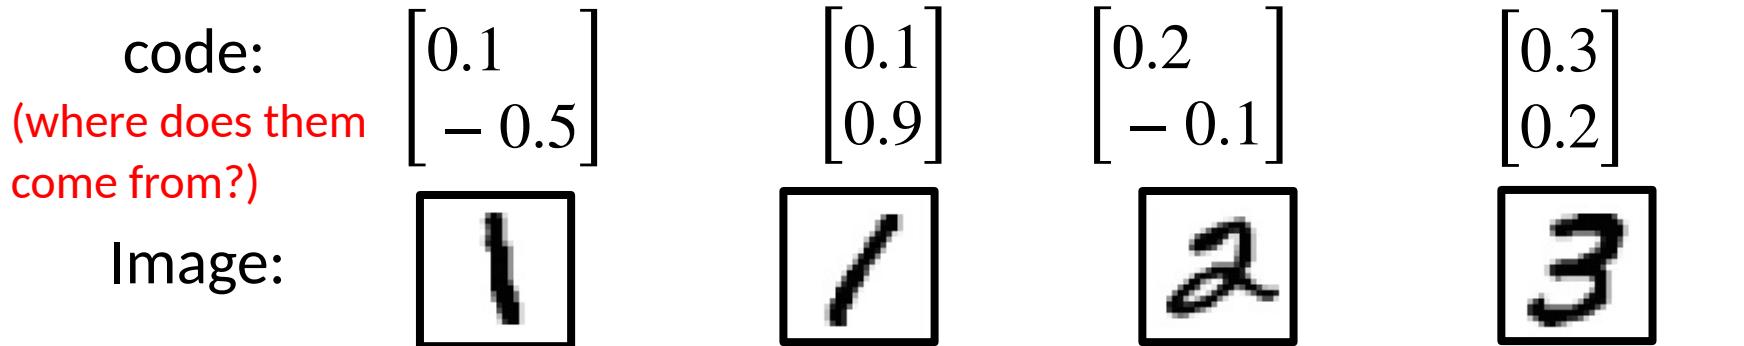
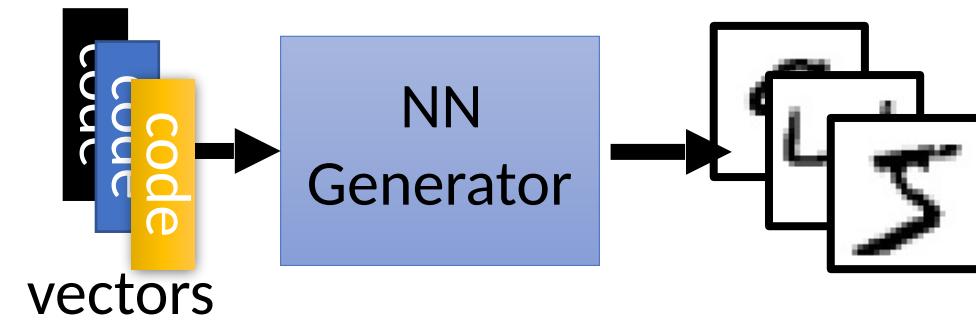
Auto-encoder

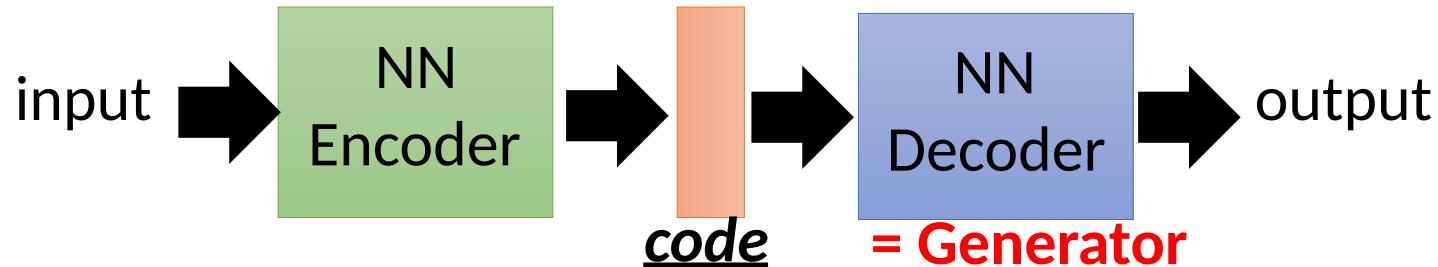


Auto-encoder

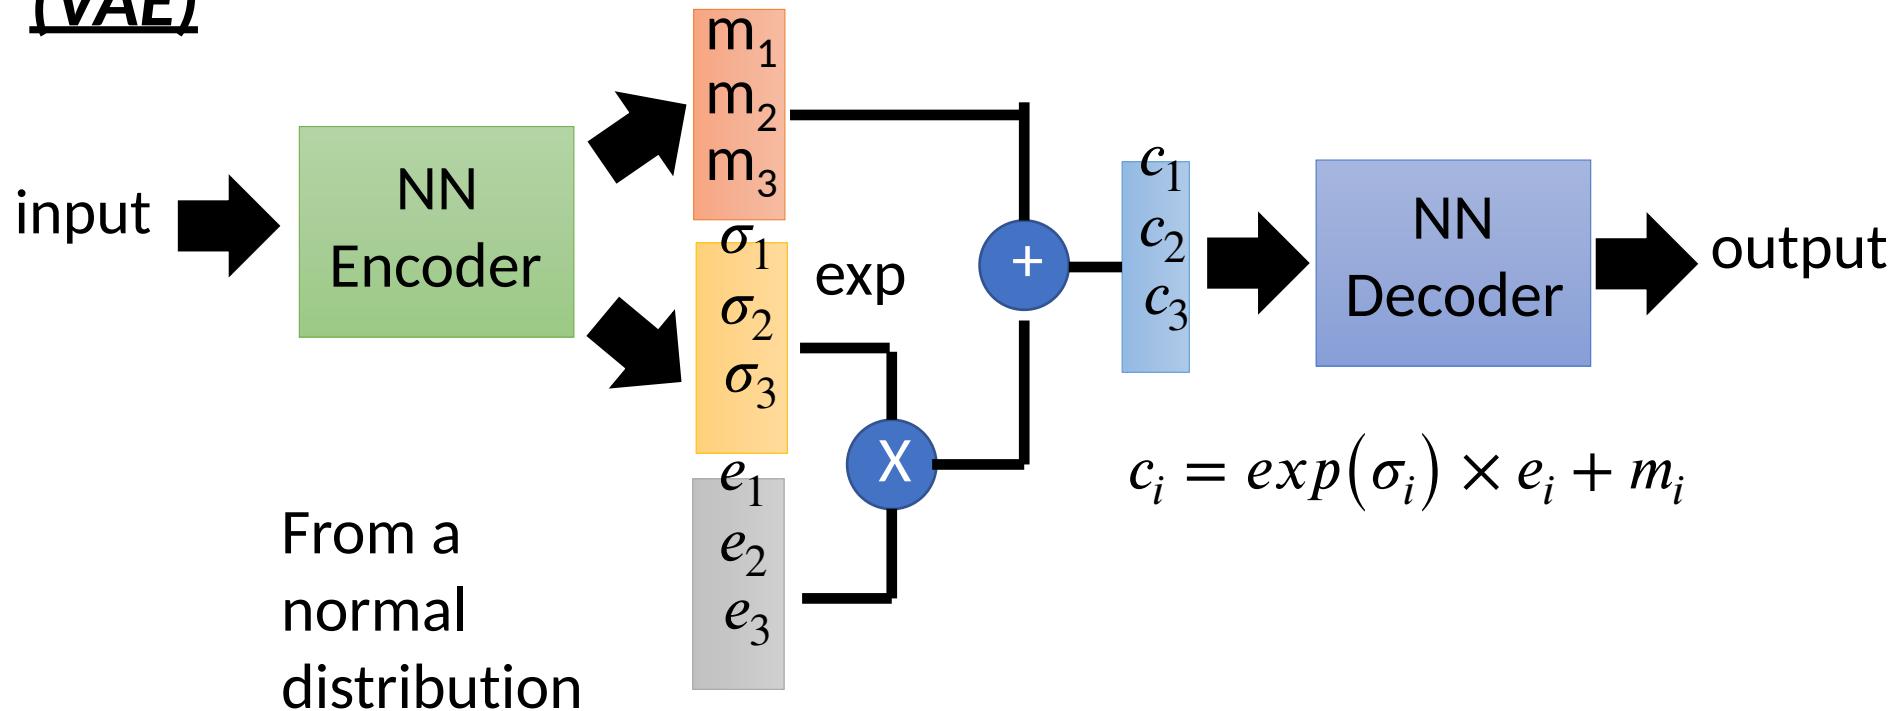


Auto-encoder

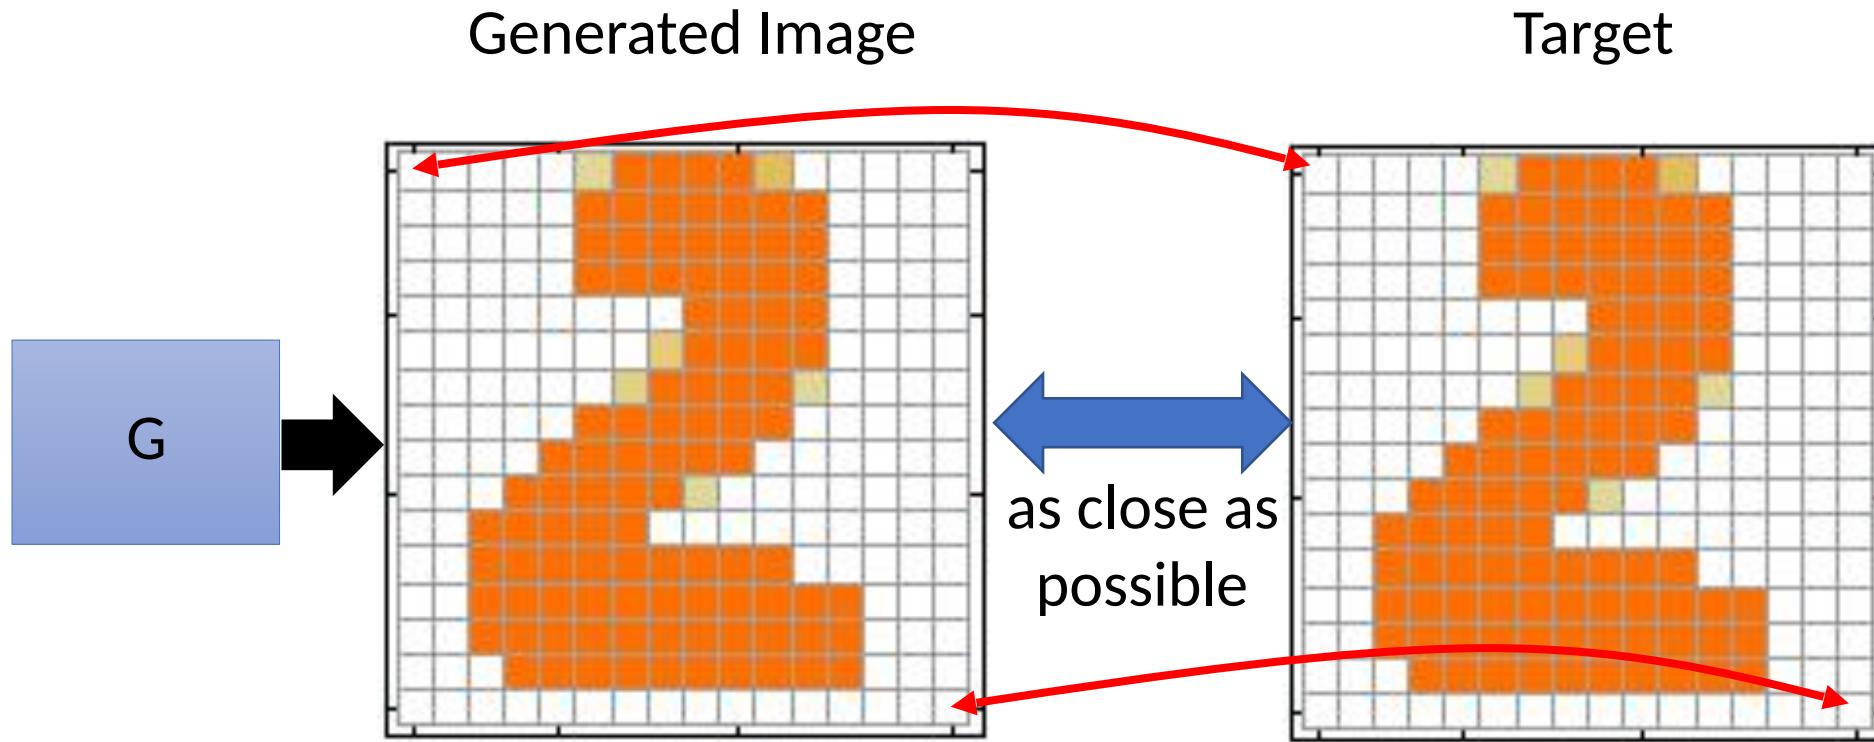




Variational Auto-encoder (VAE)



What do we miss?



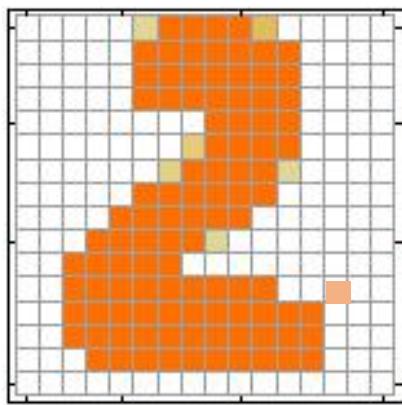
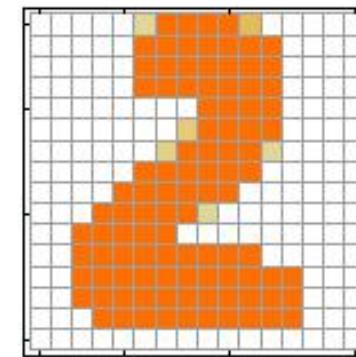
It will be fine if the generator can truly copy the target image.

What if the generator makes some mistakes

Some mistakes are serious, while some are fine.

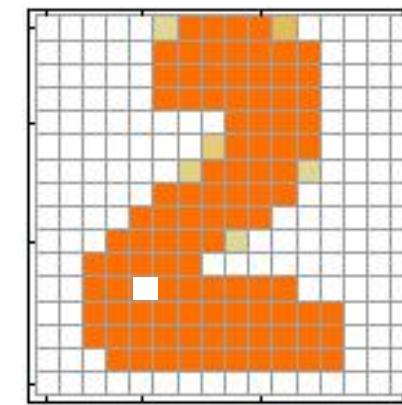
What do we miss?

Target



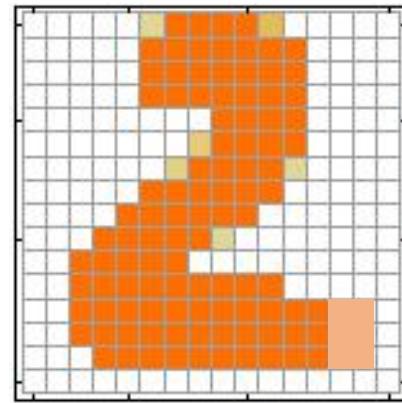
1 pixel error

我覺得不行



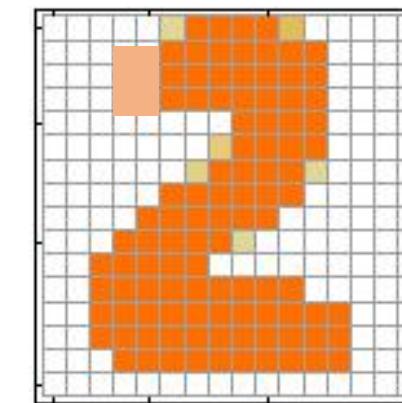
1 pixel error

我覺得不行



6 pixel errors

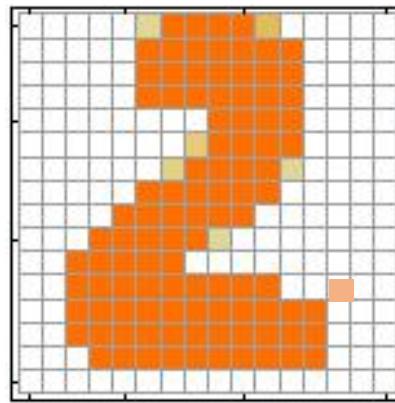
我覺得其實
可以



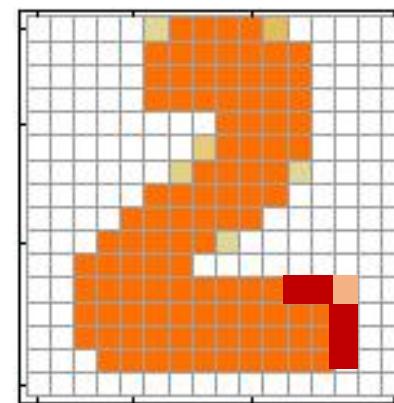
6 pixel errors

我覺得其實
可以

What do we miss?

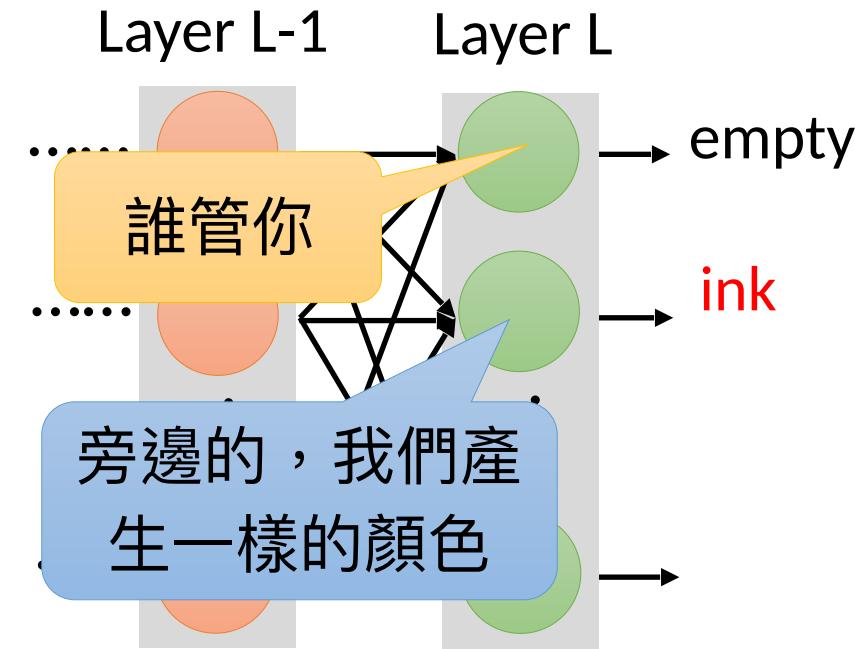


我覺得不行



我覺得其實可以

Each neural in output layer corresponds to a pixel.

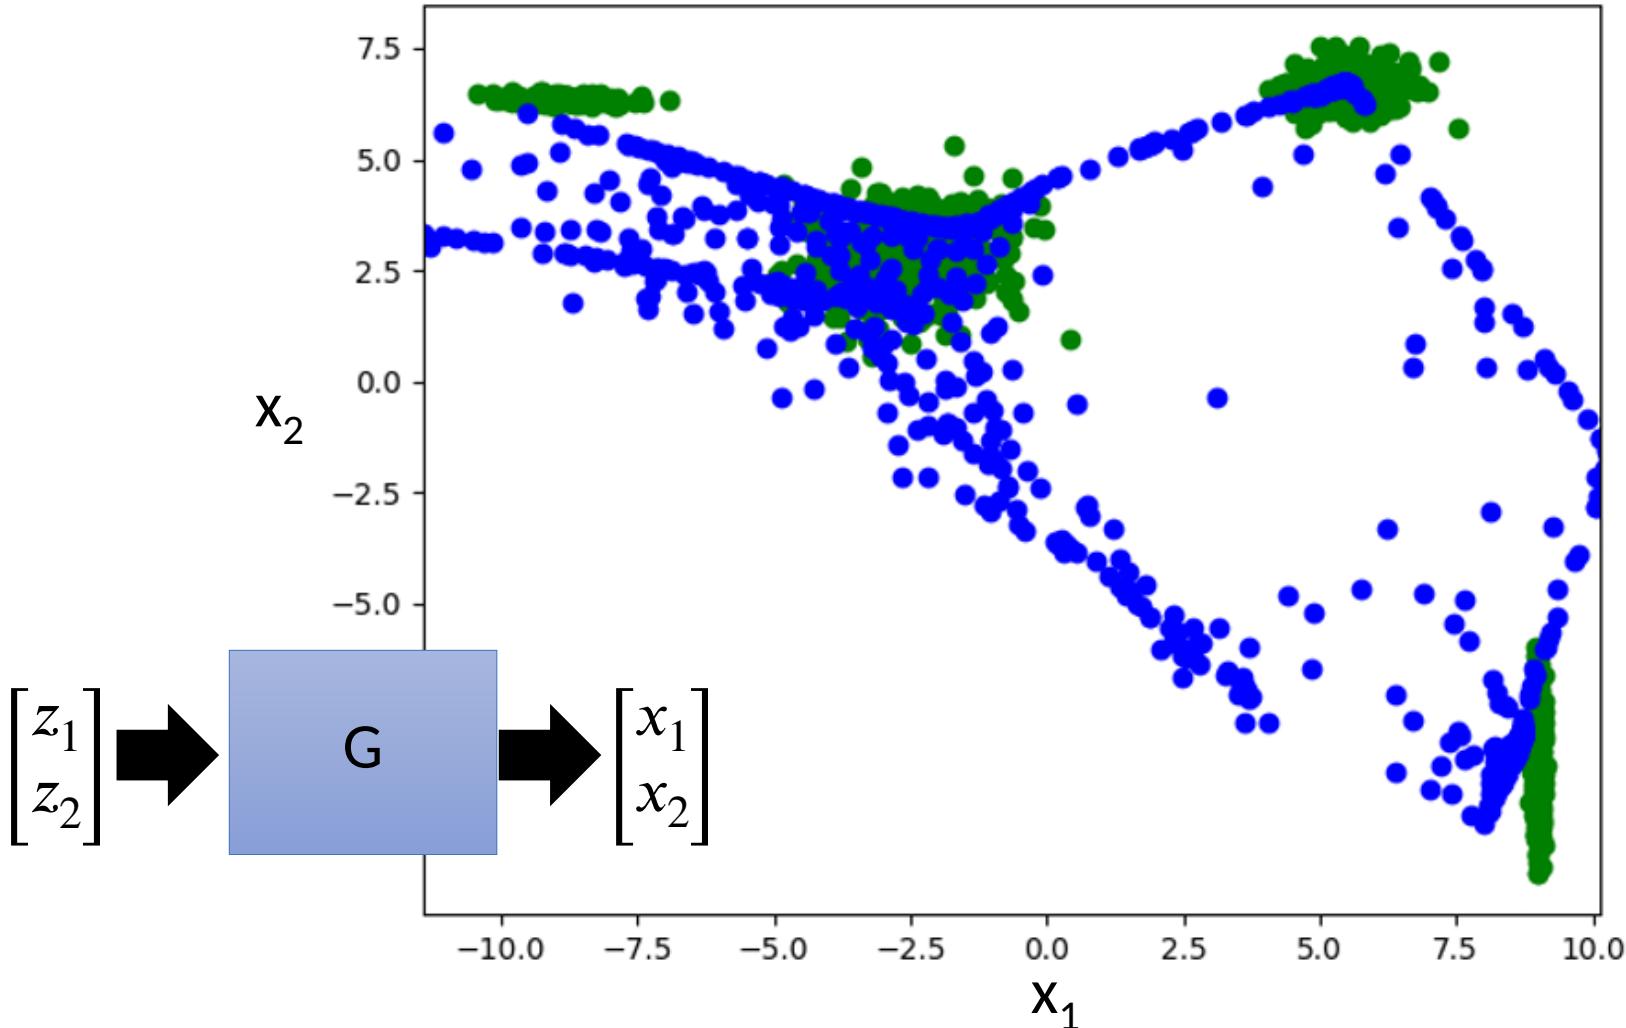


The relation between the components are critical.

Although highly correlated, they cannot influence each other.

Need deep structure to catch the relation between components.

(Variational) Auto-encoder



Can discriminator generate?

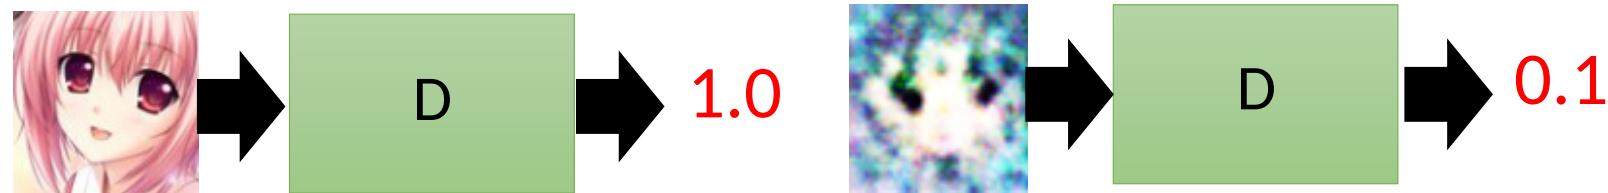
Discriminator

Evaluation function, Potential Function, Energy Function ...

- Discriminator is a function D (network, can deep)

$$D: X \rightarrow \mathbb{R}$$

- Input x : an object x (e.g. an image)
- Output $D(x)$: scalar which represents how “good” an object x is

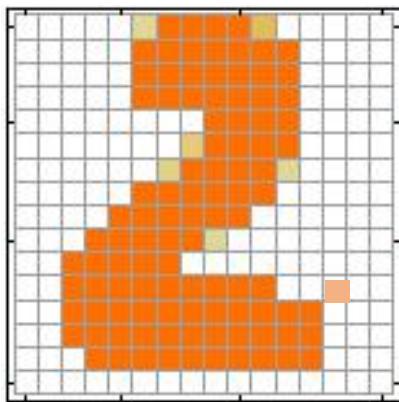


Can we use the discriminator to generate objects?

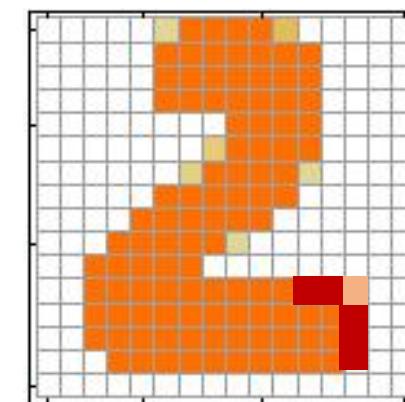
Yes.

Discriminator

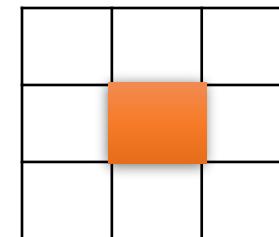
- It is easier to catch the relation between the components by top-down evaluation.



我覺得不行



我覺得其實 OK



This CNN filter is
good enough.

Discriminator

- Suppose we already have a good discriminator $D(x)$

...

Inference

- Generate object \tilde{x} that

$$\tilde{x} = \arg \max_{x \in X} D(x)$$

Enumerate all possible

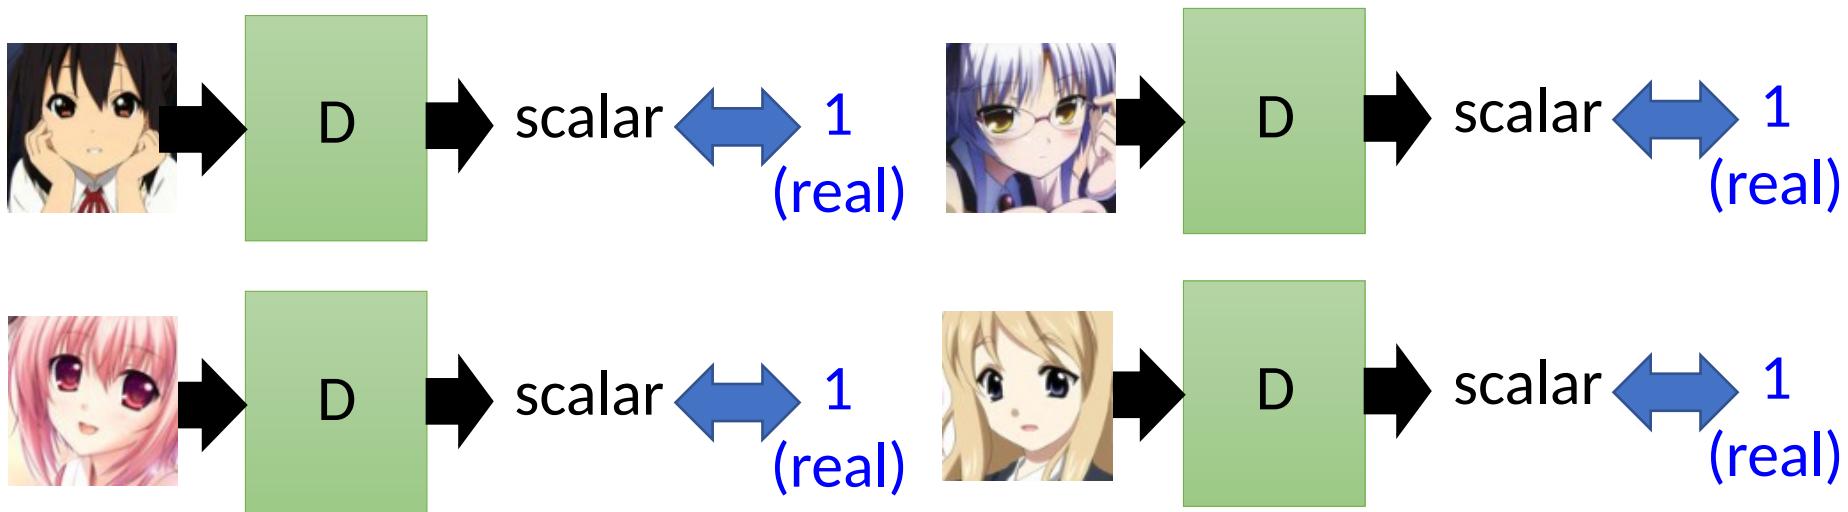
x !!!

It is feasible ???

How to learn the discriminator?

Discriminator - Training

- I have some real images

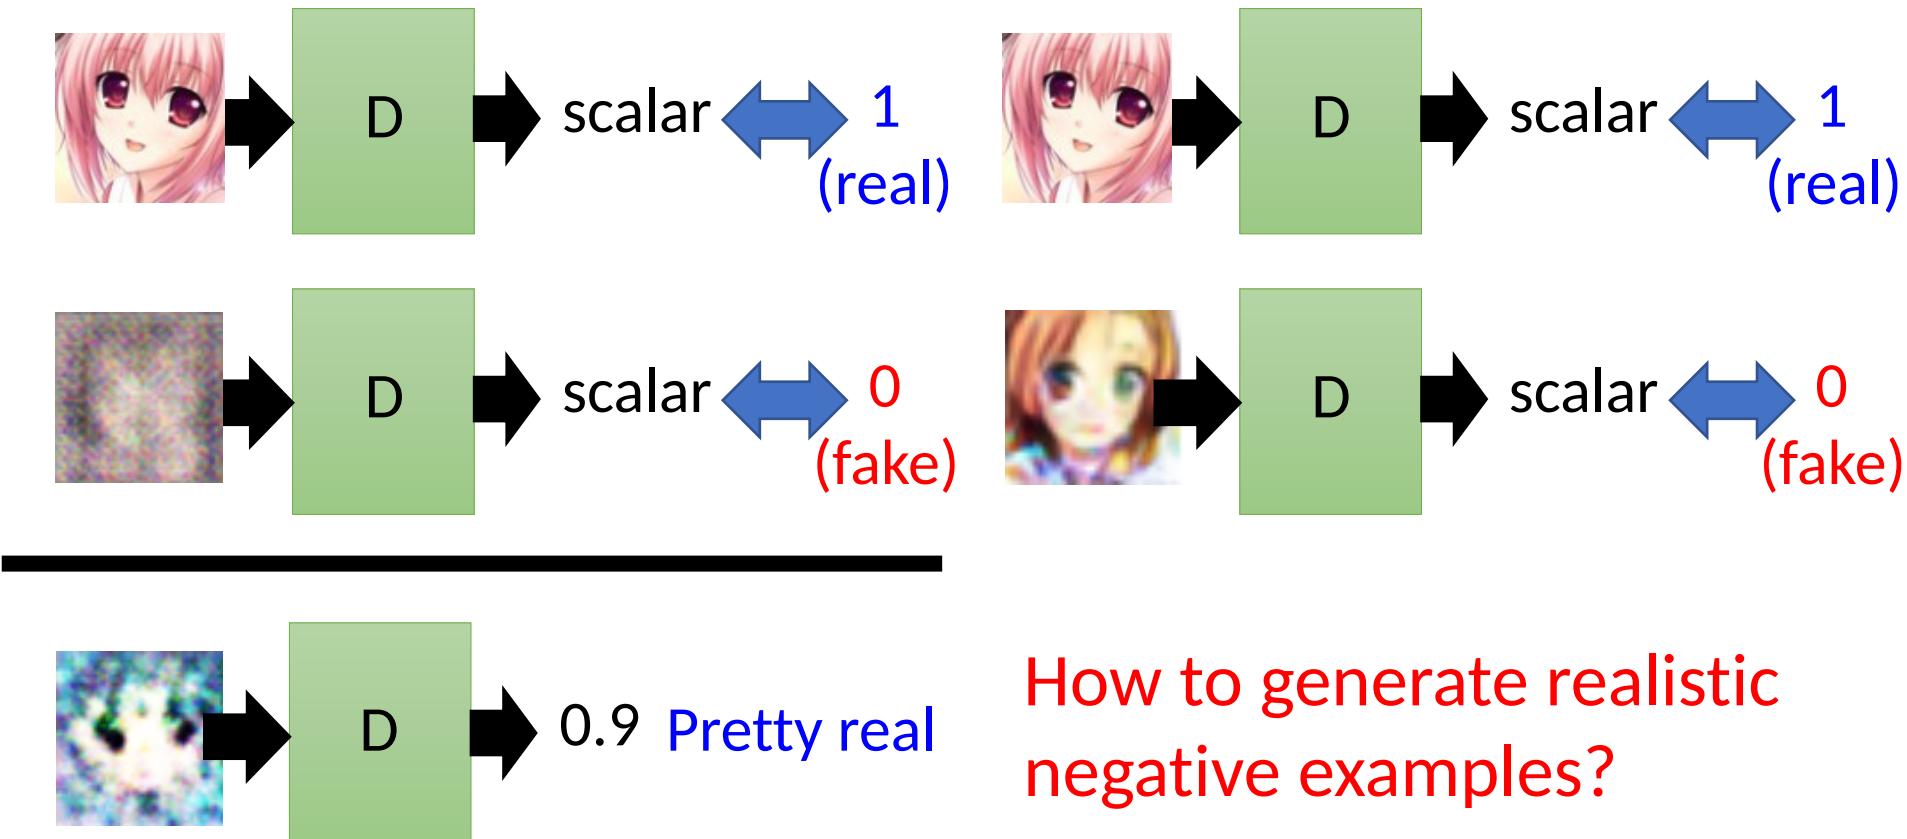


Discriminator only learns to output “1” (real).

Discriminator training needs some negative examples.

Discriminator - Training

- Negative examples are critical.



Discriminator - Training

- **General Algorithm**

- Given a set of **positive examples**, randomly generate a set of **negative examples**.



- In each iteration

- Learn a discriminator D that can discriminate positive and negative examples.



v.s.

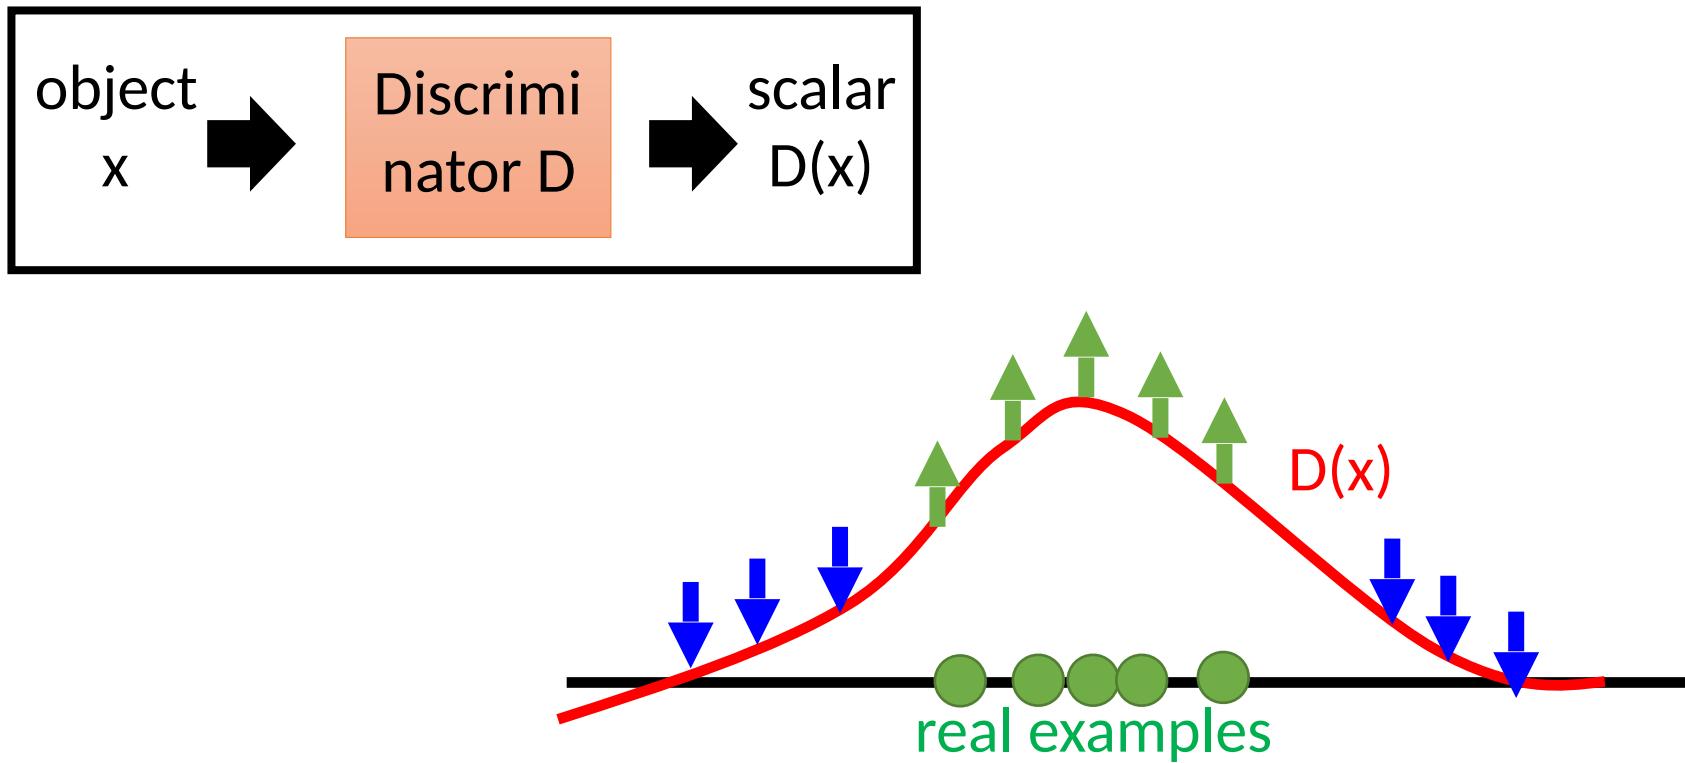


- Generate negative examples by discriminator D



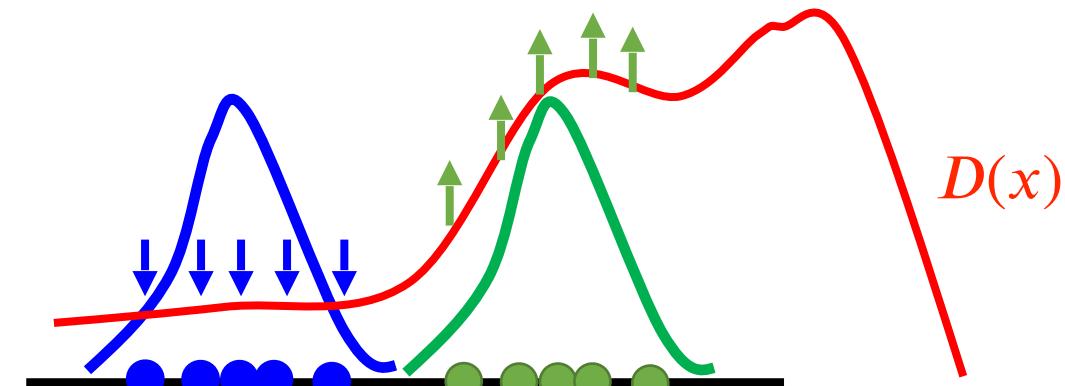
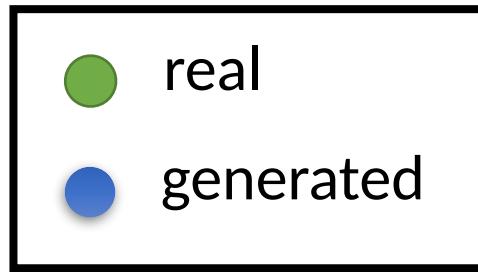
$$\tilde{x} = \arg \max_{x \in X} D(x)$$

Discriminator - Training

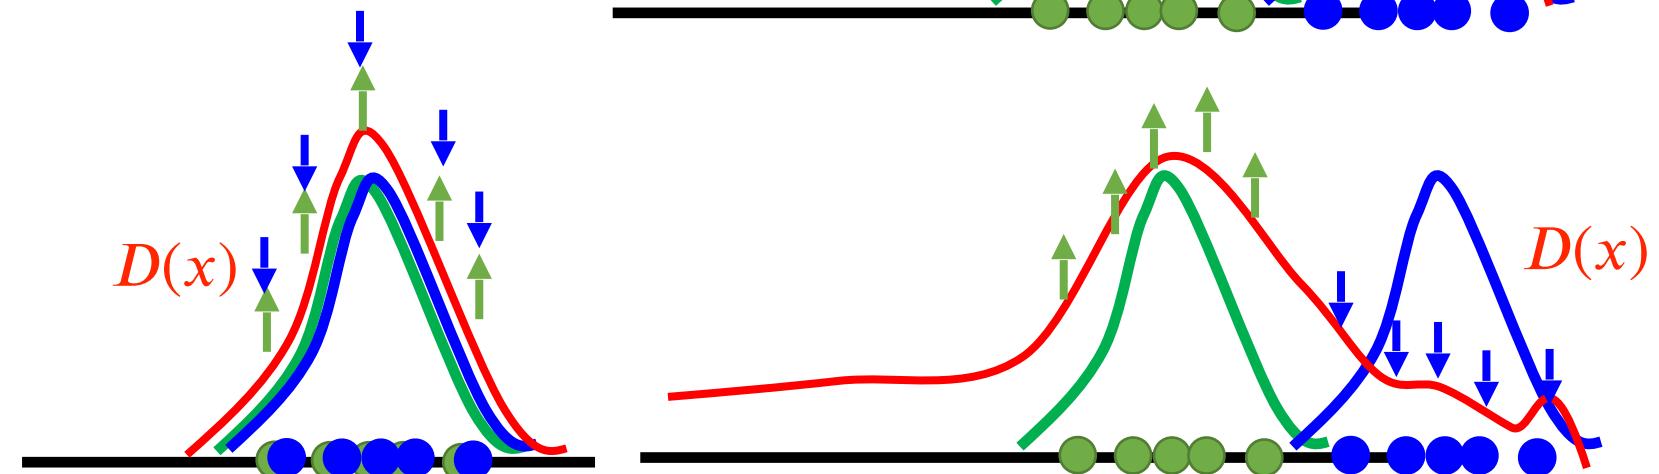
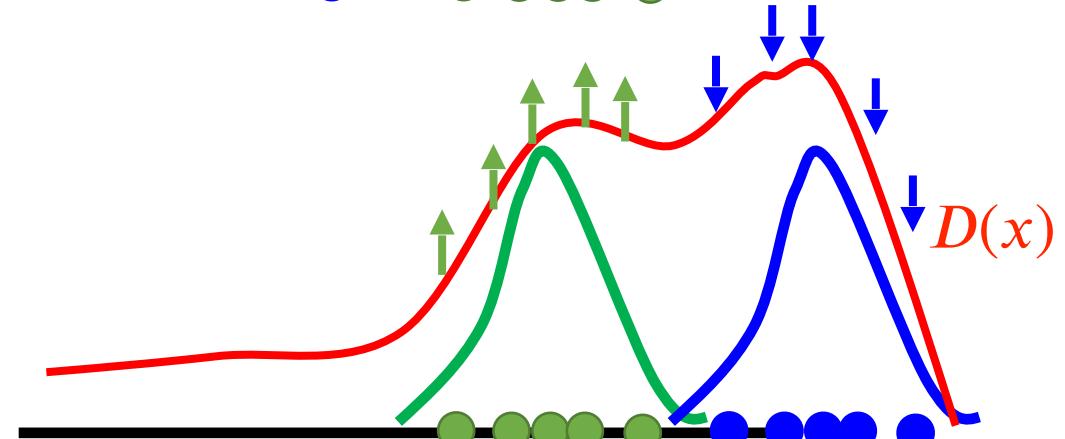


In practice, you cannot decrease all the x other than real examples.

Discriminator - Training



In the end



Generator v.s. Discriminator

- **Generator**

- Pros:

- Easy to generate even with deep model

- Cons:

- Imitate the appearance
 - Hard to learn the correlation between components

- **Discriminator**

- Pros:

- Considering the big picture

- Cons:

- Generation is not always feasible
 - Especially when your model is deep
 - How to do negative sampling?

Generator + Discriminator

- **General Algorithm**

- Given a set of **positive examples**, randomly generate a set of **negative examples**.



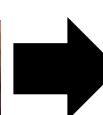
- In each iteration



- Learn a discriminator D that can discriminate positive and negative examples.



v.s.



D

- Generate negative examples by discriminator D

$$\boxed{G \rightarrow \tilde{x}}$$

$$= \boxed{\tilde{x} = \arg \max_{x \in X} D(x)}$$

Benefit of GAN

- From Discriminator's point of view
 - Using generator to generate negative samples

$$\boxed{\begin{array}{c} \text{G} \\ \longrightarrow \\ \tilde{x} \end{array}} = \boxed{\tilde{x} = \arg \max_{x \in X} D(x)}$$

efficient

- From Generator's point of view
 - Still generate the object component-by-component
 - But it is learned from the discriminator with global view.

Conditional Gan

Text-to-Image

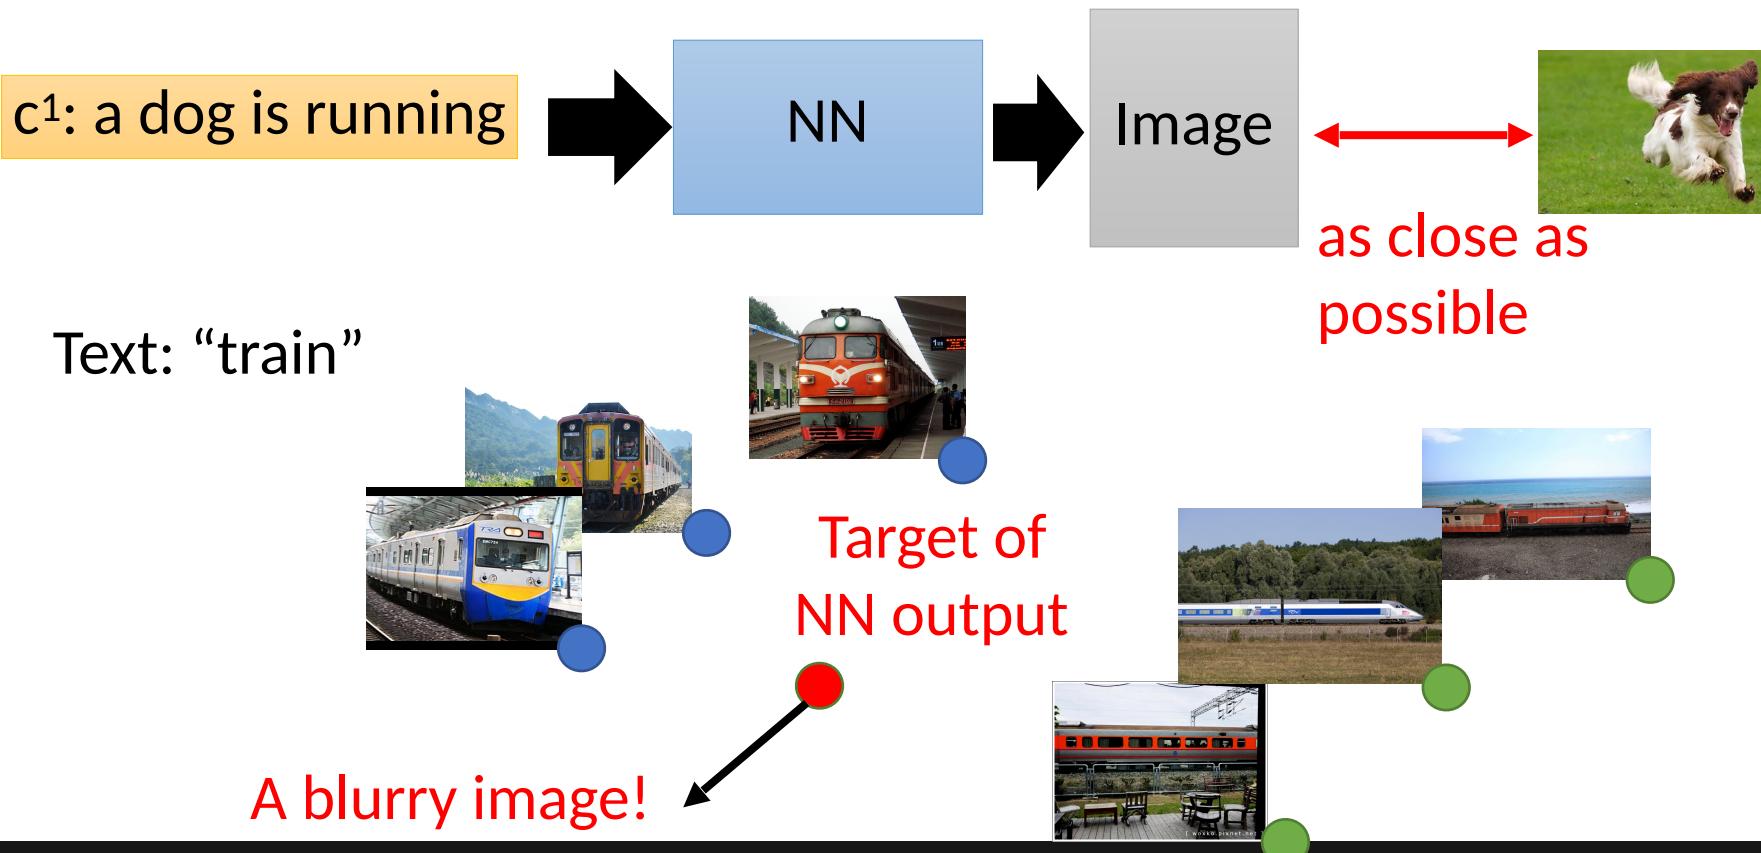
a dog is running



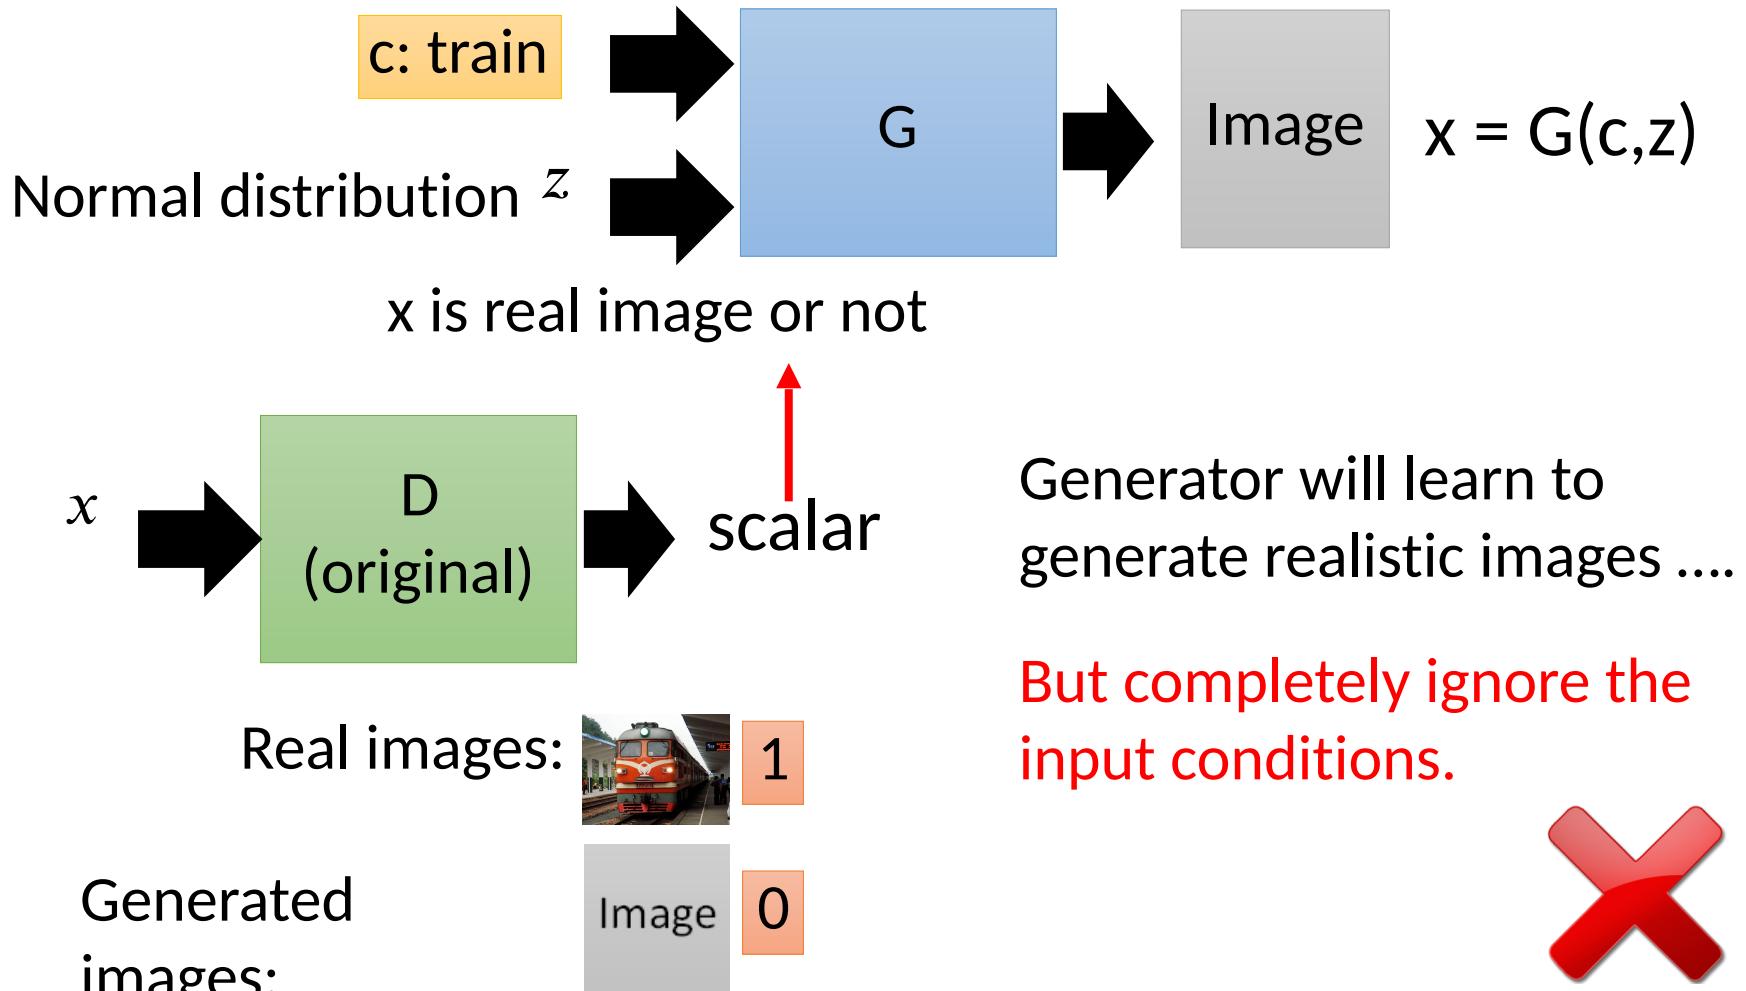
a bird is flying



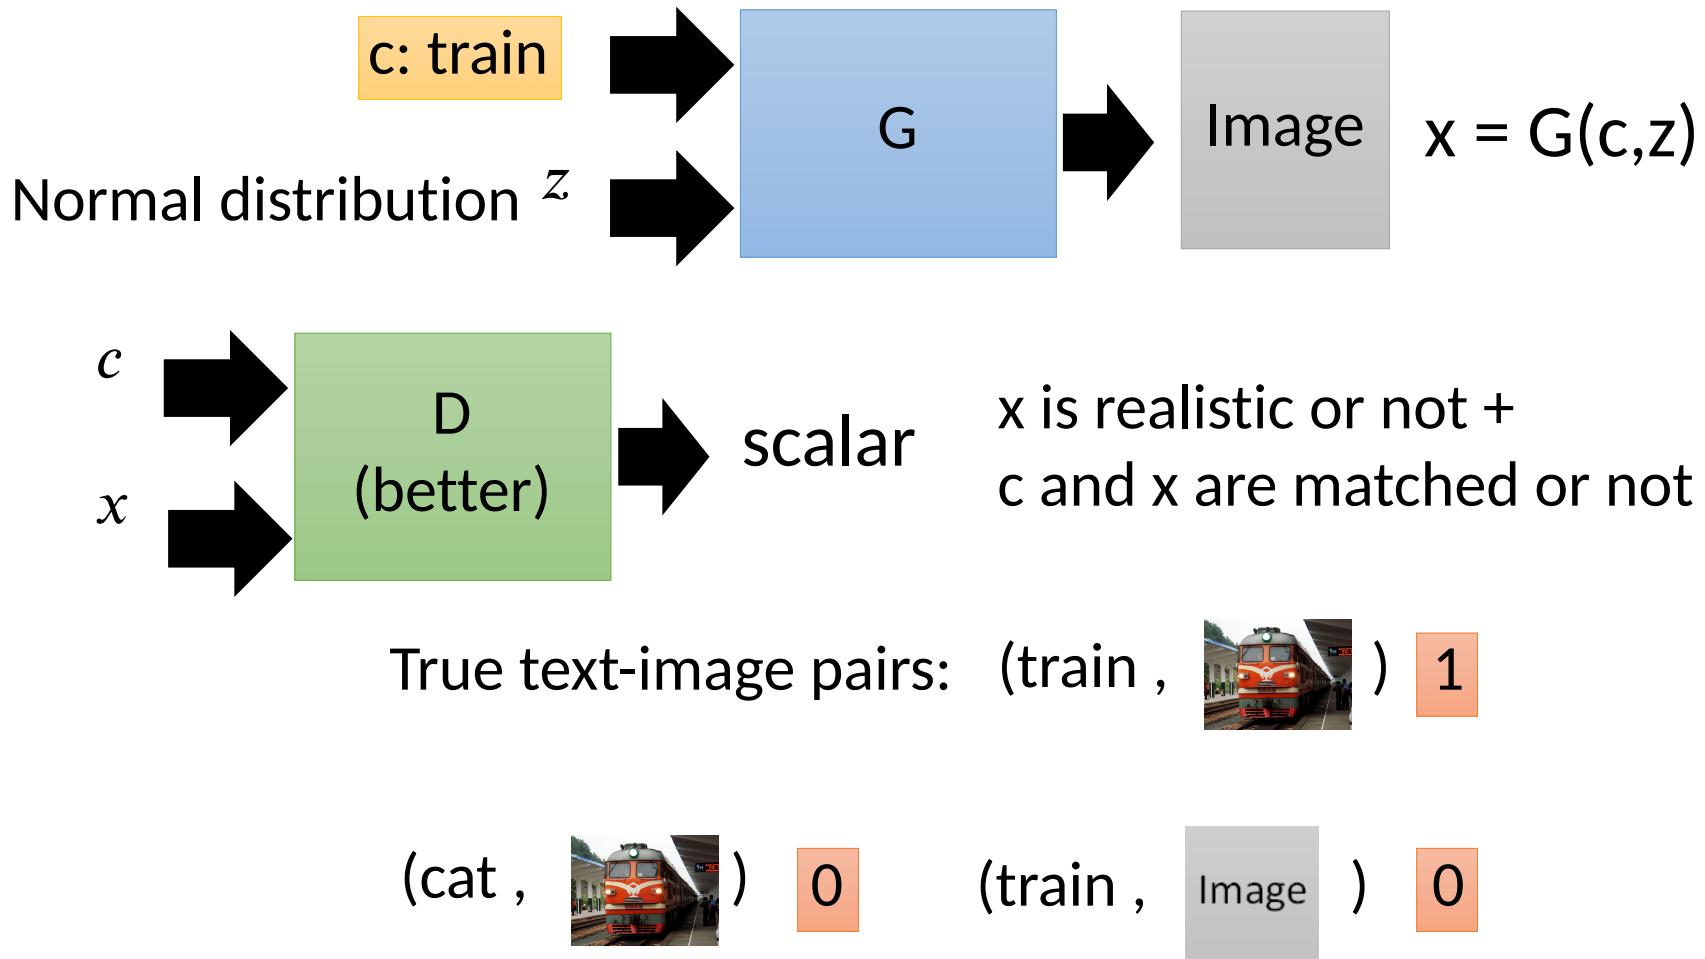
- Traditional supervised approach



Conditional GAN

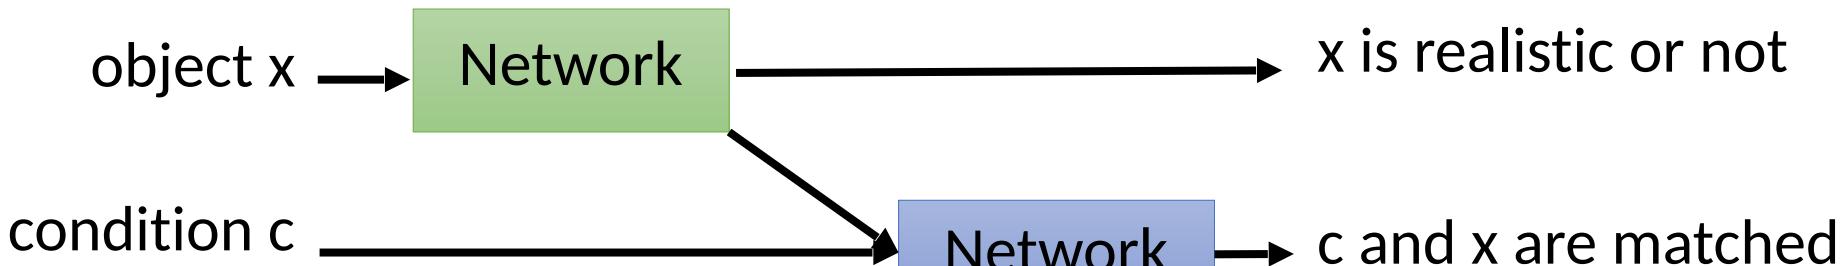
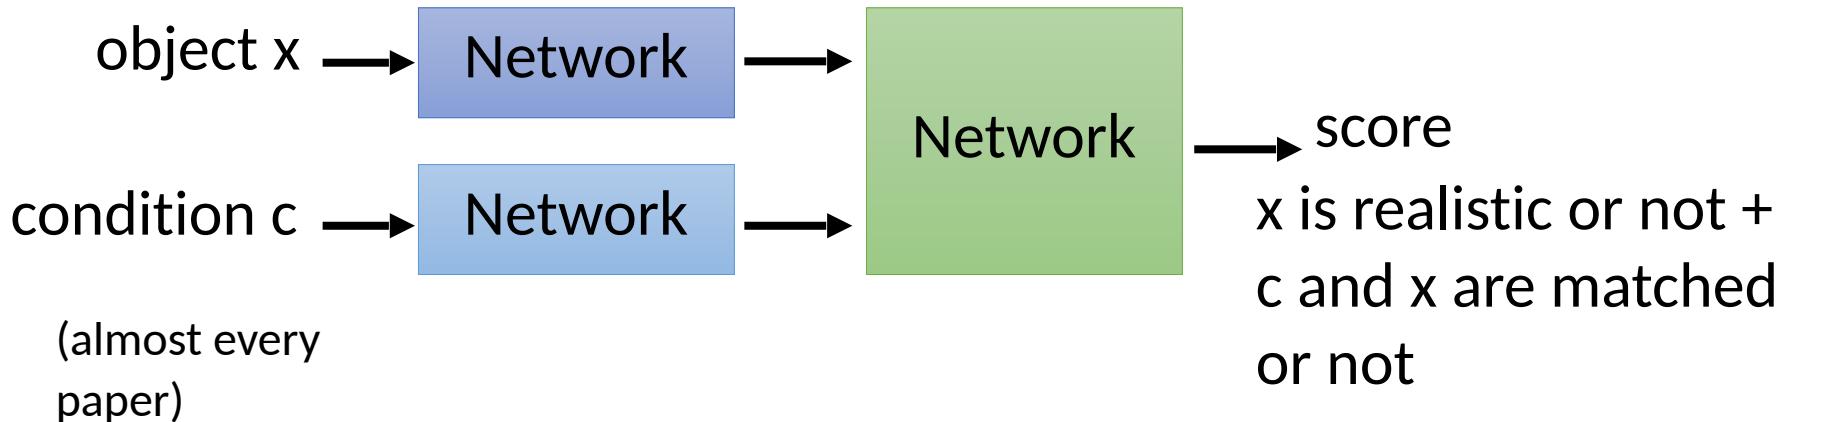


Conditional GAN



- In each training iteration:
 - Sample m positive examples $\{(c^1, x^1), (c^2, x^2), \dots, (c^m, x^m)\}$ from database
 - Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from a distribution
 - Obtaining generated data $\{\tilde{x}^1, \tilde{x}^2, \dots, \tilde{x}^m\}$, $\tilde{x}^i = G(c^i, z^i)$
 - Sample m objects $\{\hat{x}^1, \hat{x}^2, \dots, \hat{x}^m\}$ from database
 - Update discriminator parameters θ_d to maximize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log D(c^i, x^i)$
 - + $\frac{1}{m} \sum_{i=1}^m \log (1 - D(c^i, \tilde{x}^i)) + \frac{1}{m} \sum_{i=1}^m \log (1 - D(c^i, \hat{x}^i))$
 - $\theta_d \leftarrow \theta_d + \eta \nabla \tilde{V}(\theta_d)$
 - Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from a distribution
 - Sample m conditions $\{c^1, c^2, \dots, c^m\}$ from a database
 - Update generator parameters θ_g to maximize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log \left(D(G(c^i, z^i)) \right)$, $\theta_g \leftarrow \theta_g - \eta \nabla \tilde{V}(\theta_g)$

Conditional GAN - Discriminator



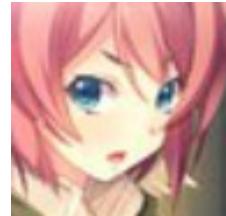
[Augustus Odena et al., ICML, 2017]

[Takeru Miyato, et al., ICLR, 2018]

[Han Zhang, et al., arXiv, 2017]

Conditional GAN

paired data



blue eyes
red hair
short hair

Collecting anime faces
and the description of its
characteristics

red hair,
green eyes



blue hair,
red eyes

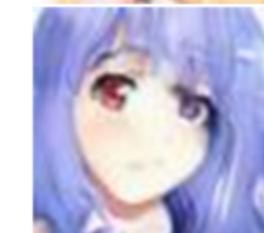
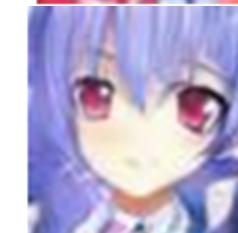
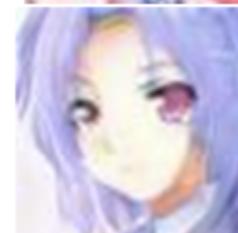
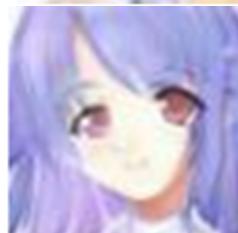


Image-to-image

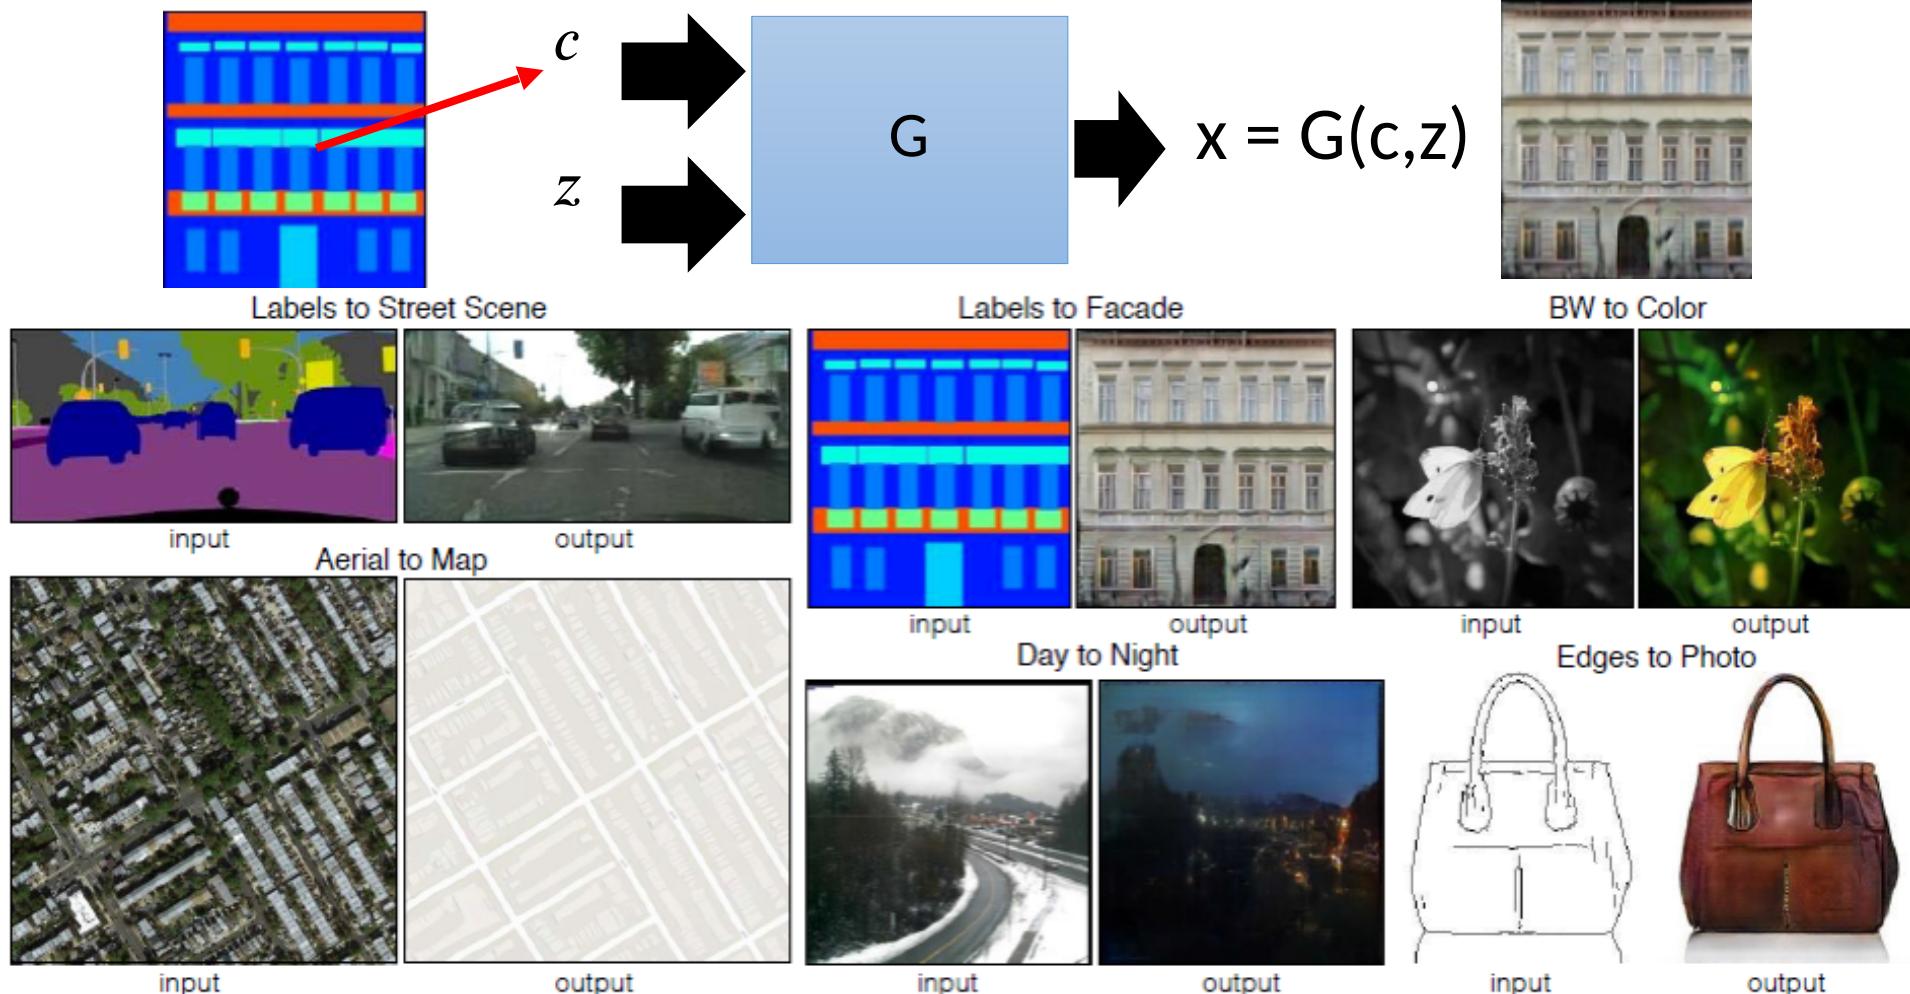
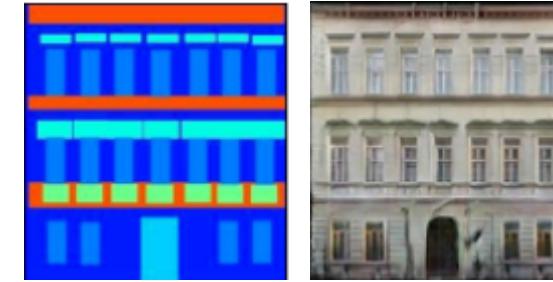
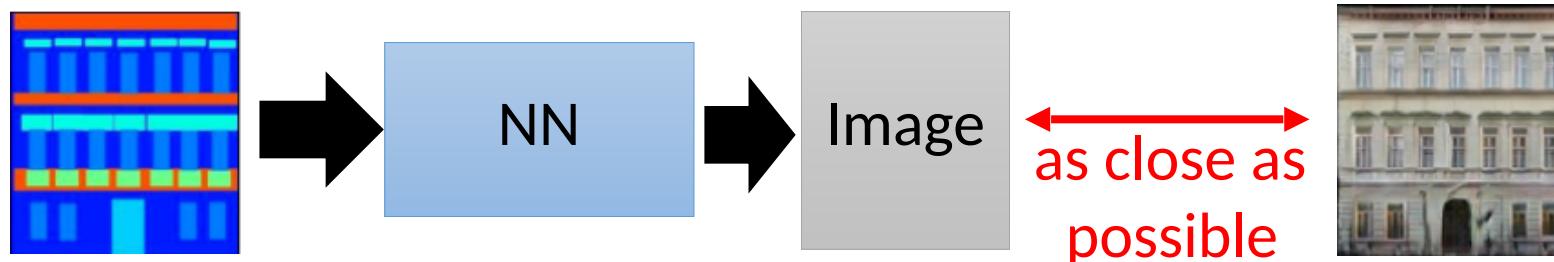


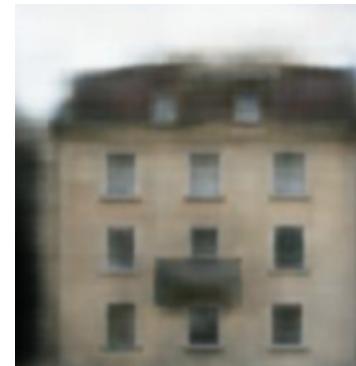
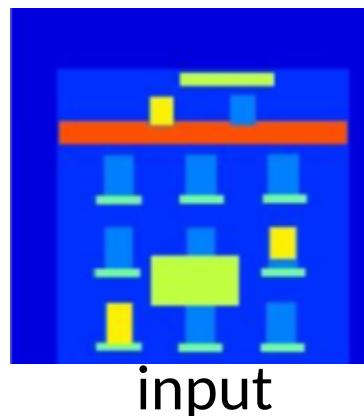
Image-to-image



- Traditional supervised approach

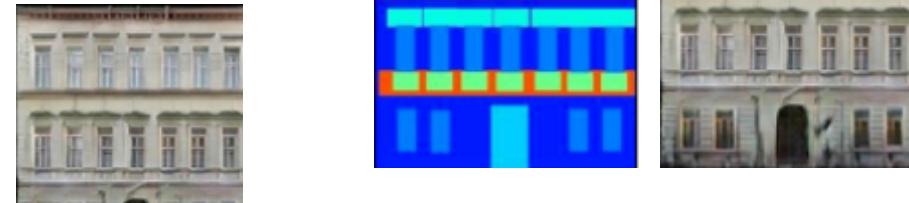


Testing:

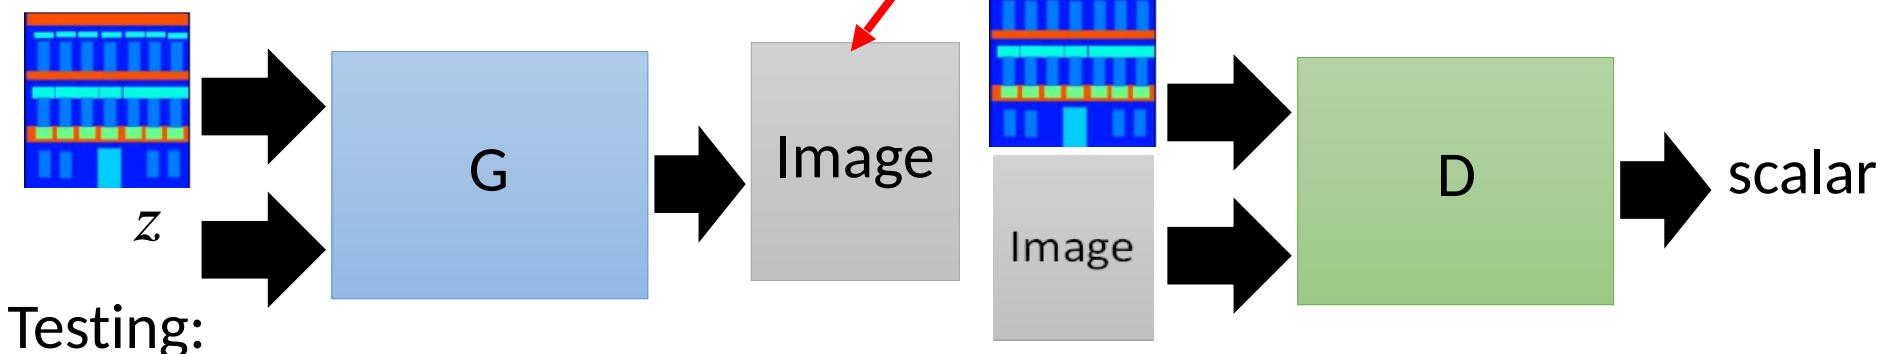


It is blurry because it is the average of several images.

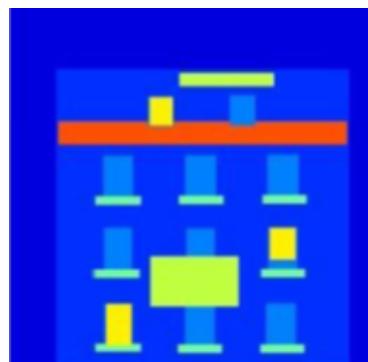
Image-to-image



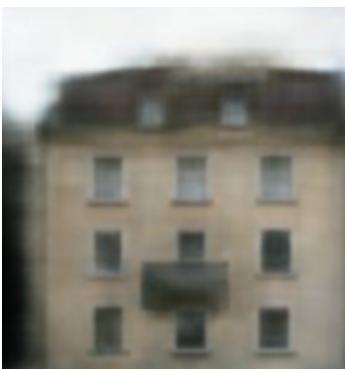
- Experimental results



Testing:



input



close



GAN



GAN + close

Conquer some issues using GANs

- Missing value imputation and GANs
- Class imbalanced dataset and GANs

Missing data imputation

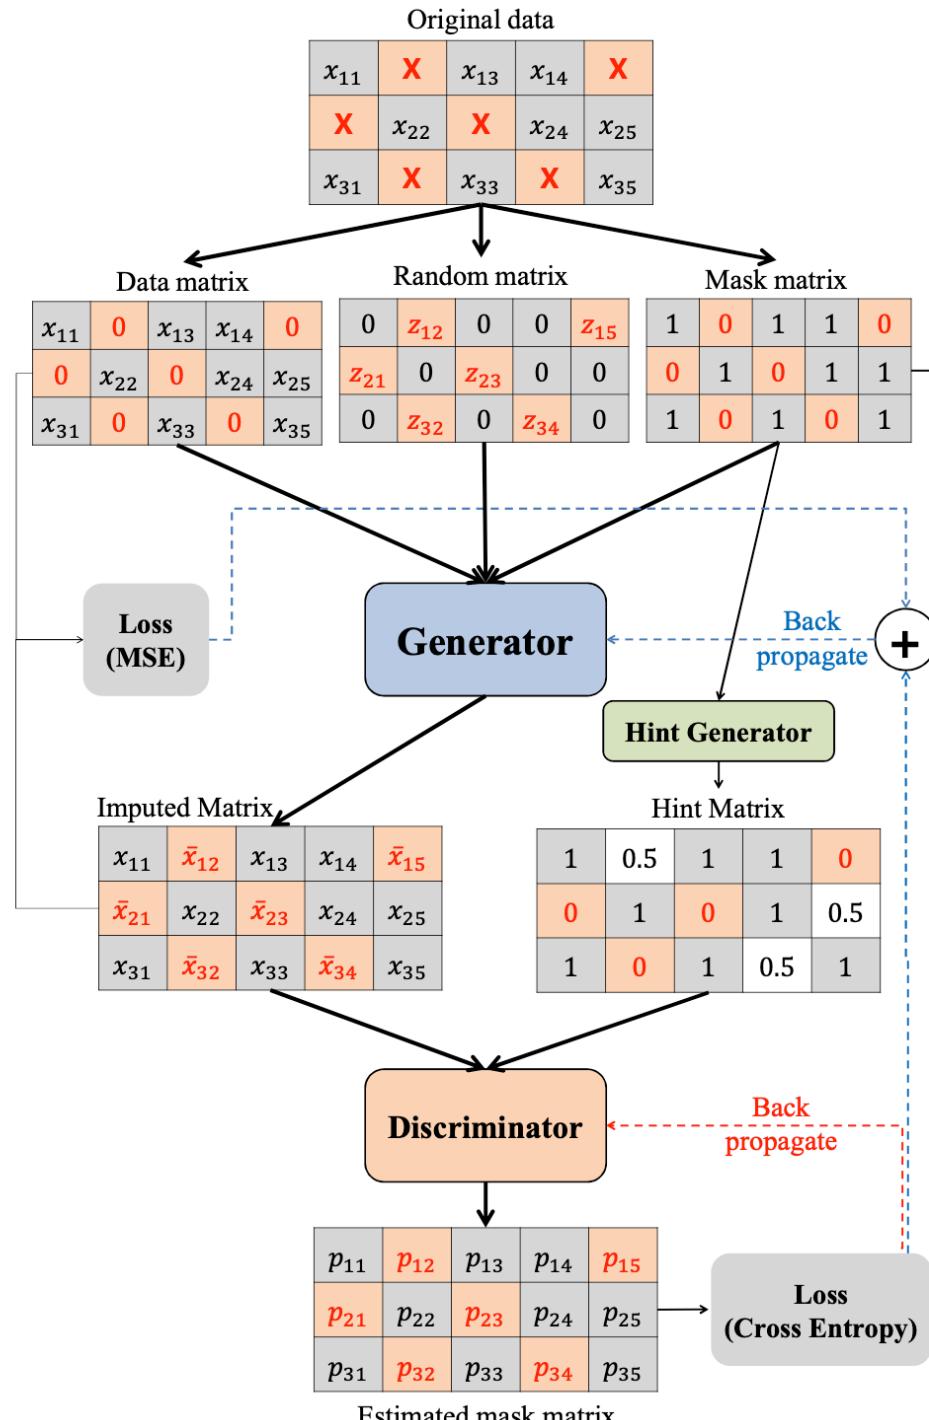
Common Ways

- Drop missing values
- Using mean/median values
- Using most frequent values
- Using k-nn
- Using multivariate imputation by chained equations (MICE)

Based on GAN

- Generative Adversarial Imputation Nets (GAIN)

The architecture of GAIN



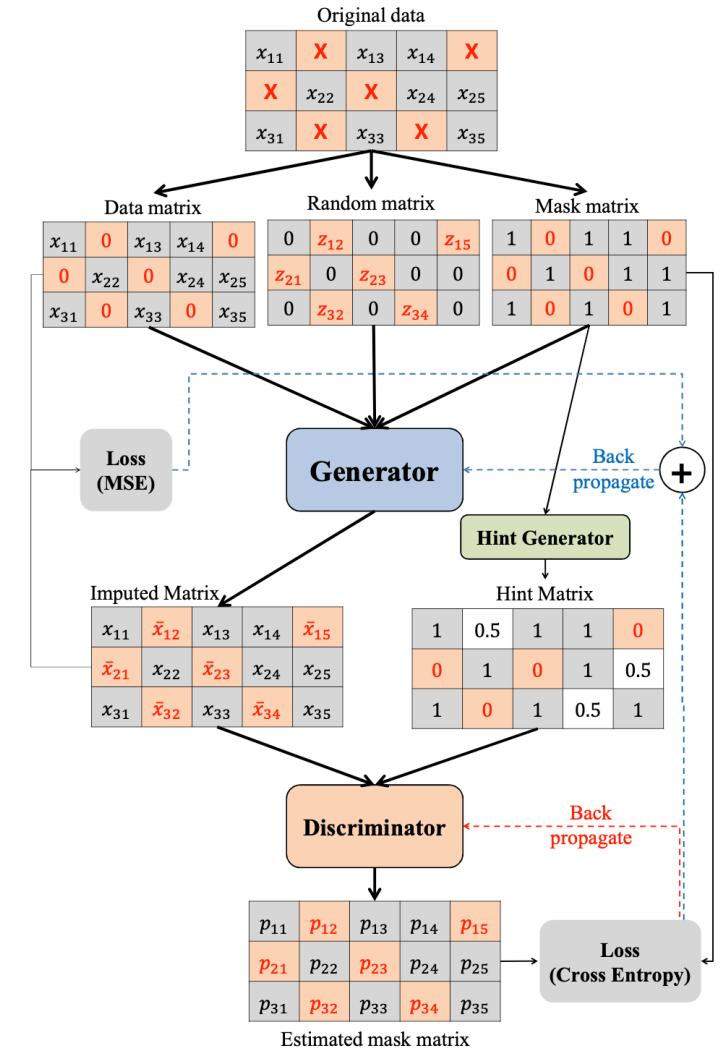
Given d-dimensional space $\mathcal{X} = \mathcal{X}_1 \times \dots \times \mathcal{X}_d$

$\mathbf{X} = (X_1, \dots, X_d)$ is a random variable taking values in X

$\mathbf{M} = (M_1, \dots, M_d)$, taking values in $\{0, 1\}$

$\tilde{\mathbf{X}} = (\tilde{X}_1, \dots, \tilde{X}_d) \in \tilde{\mathcal{X}}$

$$\tilde{X}_i = \begin{cases} X_i, & \text{if } M_i = 1 \\ *, & \text{otherwise} \end{cases}$$

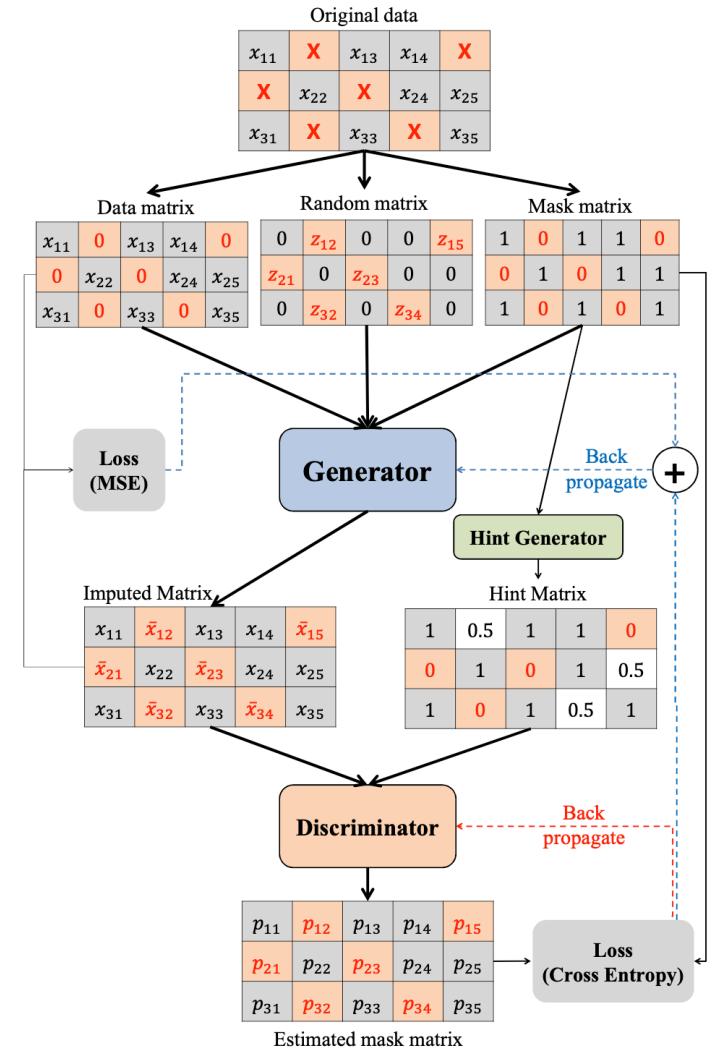


$$\bar{\mathbf{X}} = G(\tilde{\mathbf{X}}, \mathbf{M}, (\mathbf{1} - \mathbf{M}) \odot \mathbf{Z})$$

$$\hat{\mathbf{X}} = \mathbf{M} \odot \bar{\mathbf{X}} + (\mathbf{1} - \mathbf{M}) \odot \bar{\mathbf{X}}$$

Why Hint Generator?

```
H_mb_temp = binary_sampler_hint_rate, batch_size, dim)
H_mb = M_mb * H_mb_temp
```



Deal with class imbalanced dataset

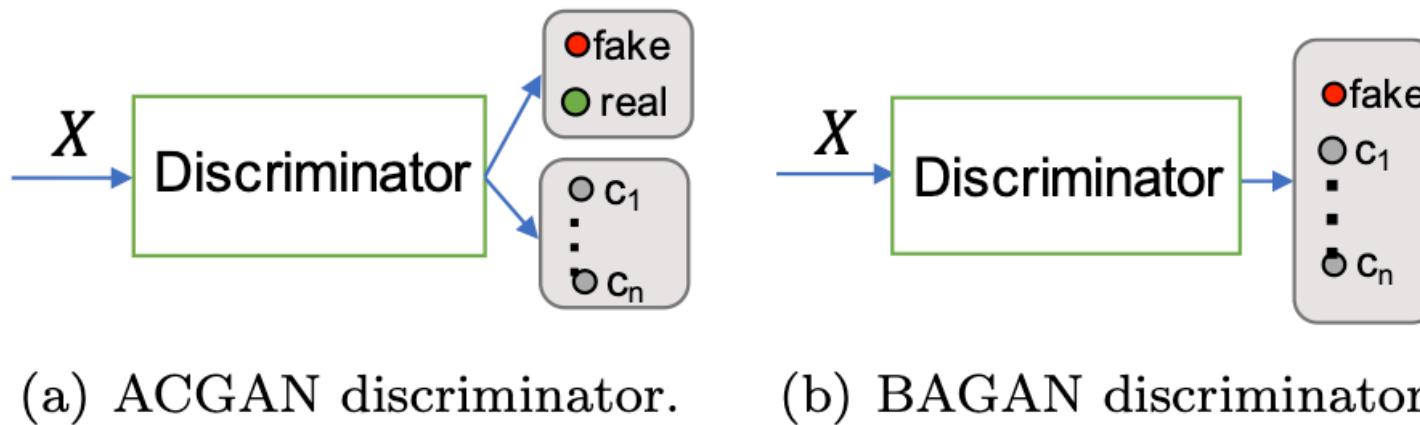
Common ways

- Over-sampling/Under-sampling
- Synthetic Minority Oversampling Technique (SMOTE)

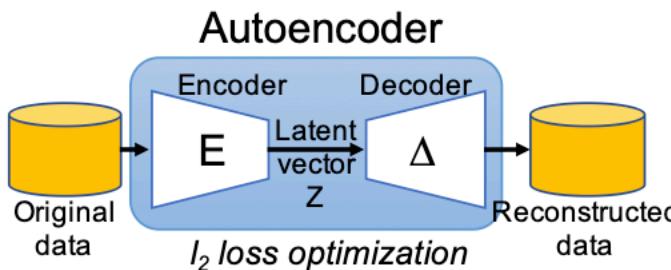
Based on GAN

- Balancing GAN (BAGAN)

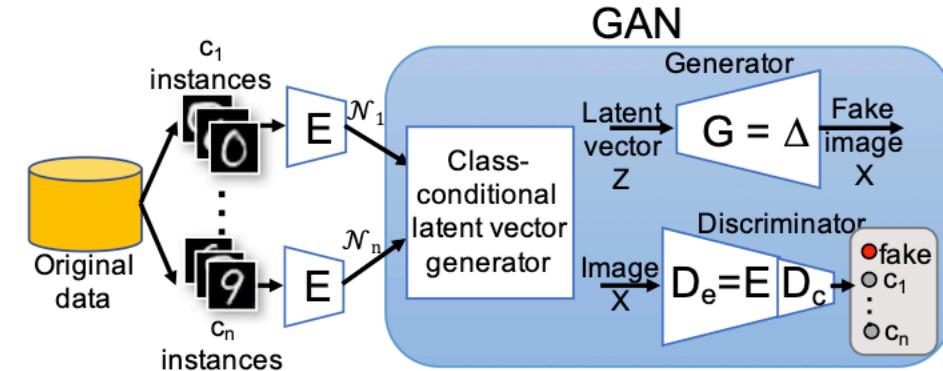
Discuss difference between two discriminators



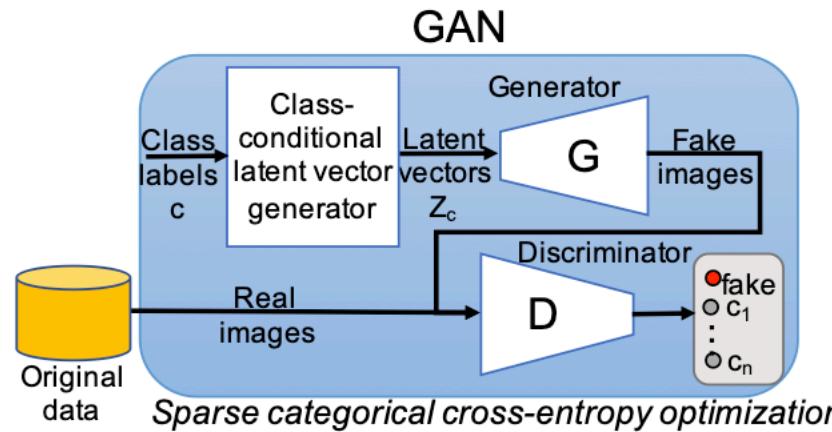
BAGAN Architecture



(a) Autoencoder training.

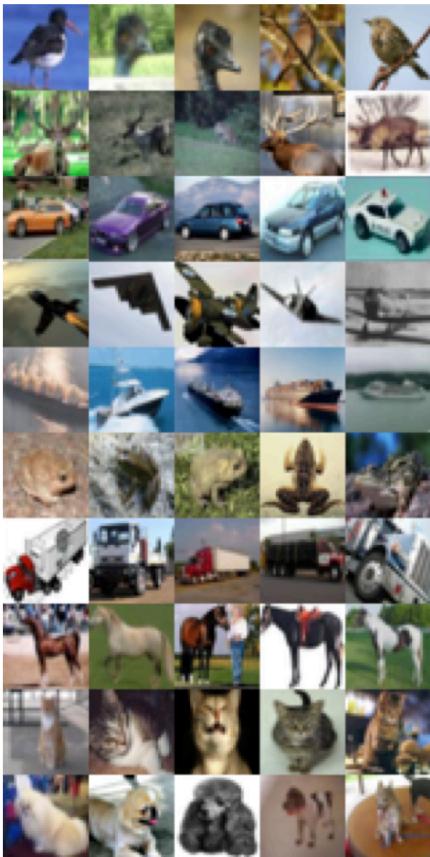


(b) GAN initialization.



(c) GAN training.

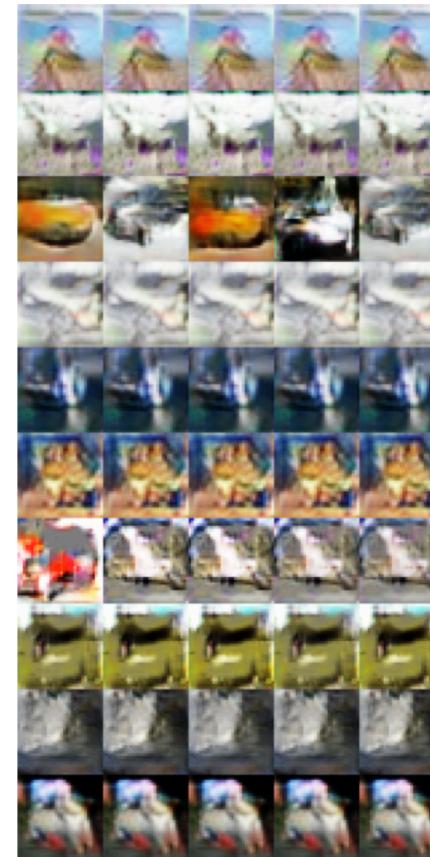
Results



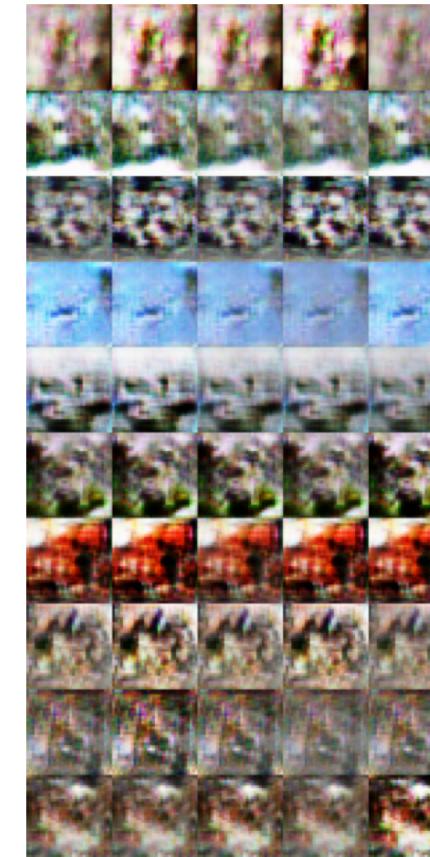
(a) Real image samples



(b) *BAGAN*



(c) *ACGAN*



(d) *Simple GAN*

Implementation with PyTorch

https://github.com/TzuYuOu/GAN_with_PyTorch

Generator

```
class Generator(nn.Module):
    def __init__(self):
        super(Generator, self).__init__()

        def block(in_feat, out_feat, normalize=True):
            layers = [nn.Linear(in_feat, out_feat)]
            if normalize:
                layers.append(nn.BatchNorm1d(out_feat, 0.8))
            layers.append(nn.LeakyReLU(0.2, inplace=True))
            return layers

        self.model = nn.Sequential(
            *block(latent_dim, 128, normalize=False),
            *block(128, 256),
            *block(256, 512),
            *block(512, 1024),
            nn.Linear(1024, int(np.prod(img_shape))),
            nn.Tanh()
        )

    def forward(self, z):
        img = self.model(z)
        img = img.view(img.size(0), *img_shape)
        return img
```

Discriminator

```
class Discriminator(nn.Module):
    def __init__(self):
        super(Discriminator, self).__init__()

        self.model = nn.Sequential(
            nn.Linear(int(np.prod(img_shape)), 512),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Linear(512, 256),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Linear(256, 1),
            nn.Sigmoid(),
        )

    def forward(self, img):
        img_flat = img.view(img.size(0), -1)
        validity = self.model(img_flat)

    return validity
```

Training preparation

```
cuda = True if torch.cuda.is_available() else False
# Loss function
adversarial_loss = torch.nn.BCELoss()

# Initialize generator and discriminator
generator = Generator()
discriminator = Discriminator()

if cuda:
    generator.cuda()
    discriminator.cuda()
    adversarial_loss.cuda()

# Configure data loader
os.makedirs("../..../data/mnist", exist_ok=True)
dataloader = torch.utils.data.DataLoader(
    datasets.MNIST(
        "../..../data/mnist",
        train=True,
        download=True,
        transform=transforms.Compose(
            [transforms.Resize(img_size), transforms.ToTensor(), transforms.Normalize([0.5], [0.5])]
        ),
    ),
    batch_size=batch_size,
    shuffle=True,
)
Tensor = torch.cuda.FloatTensor if cuda else torch.FloatTensor

# Optimizers
optimizer_G = torch.optim.Adam(generator.parameters(), lr=lr, betas=(b1,b2))
optimizer_D = torch.optim.Adam(discriminator.parameters(), lr=lr, betas=(b1, b2))
```

Training

```
for epoch in range(n_epochs):
    for i, (imgs, _) in enumerate(dataloader):

        # Adversarial ground truths
        valid = Variable(Tensor(imgs.size(0), 1).fill_(1.0), requires_grad=False)
        fake = Variable(Tensor(imgs.size(0), 1).fill_(0.0), requires_grad=False)

        # Configure input
        real_imgs = Variable(imgs.type(Tensor))

        # -----
        # Train Generator
        # -----

        optimizer_G.zero_grad()

        # Sample noise as generator input
        z = Variable(Tensor(np.random.normal(0, 1, (imgs.shape[0], latent_dim)))))

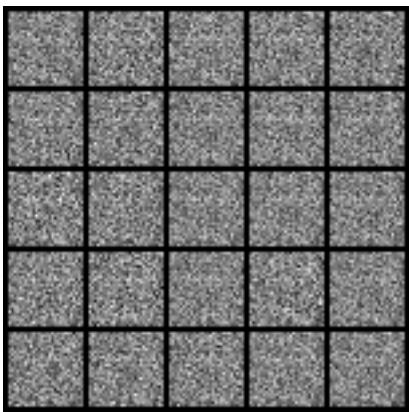
        # Generate a batch of images
        gen_imgs = generator(z)

        # Loss measures generator's ability to fool the discriminator
        g_loss = adversarial_loss(discriminator(gen_imgs), valid)

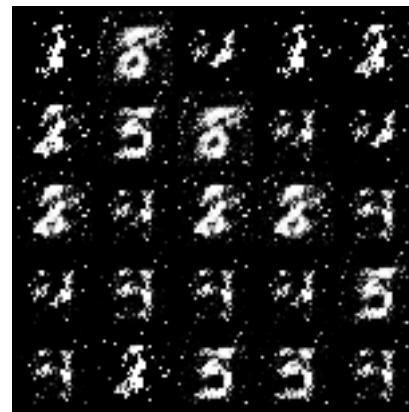
        g_loss.backward()
        optimizer_G.step()
```

```
# -----
# Train Discriminator
# -----  
  
optimizer_D.zero_grad()  
  
# Measure discriminator's ability to classify real from generated samples
real_loss = adversarial_loss(discriminator(real_imgs), valid)
fake_loss = adversarial_loss(discriminator(gen_imgs.detach()), fake)
d_loss = (real_loss + fake_loss) / 2  
  
d_loss.backward()
optimizer_D.step()  
  
print(
    "[Epoch %d/%d] [Batch %d/%d] [D loss: %f] [G loss: %f]"
    % (epoch, n_epochs, i, len(dataloader), d_loss.item(), g_loss.item())
)  
  
batches_done = epoch * len(dataloader) + i
if batches_done % sample_interval == 0:
    save_image(gen_imgs.data[:25], "images/%d.png" % batches_done, nrow=5, normalize=True)
```

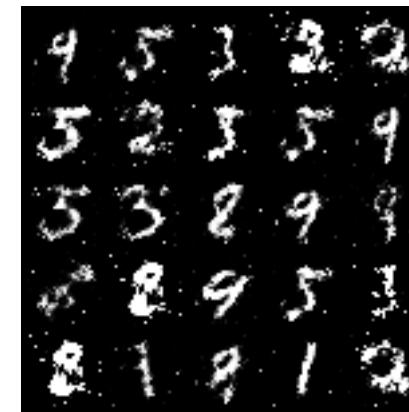
Results



epoch1



epoch20



epoch50

Results



epoch100



epoch150



epoch200

Reference

- Machine Learning and having it deep and structured
(2018, Spring) http://speech.ee.ntu.edu.tw/~tlkagk/courses_MLDS18.html
- Missing Data Imputation using Generative Adversarial Nets
<https://arxiv.org/abs/1806.02920>
- Data Argumentation with Balancing GAN <https://arxiv.org/abs/1803.09655>
- PyTorch-GAN <https://github.com/eriklindernoren/PyTorch-GAN>
- ACGAN <https://arxiv.org/abs/1610.09585>